

AN ARCHIVAL ALMA SEARCH FOR UNRESOLVED SPECTRAL-CARRIER TECHNOSIGNATURES TOWARD 82 STELLAR SYSTEMS WITHIN 40 PC AT 90–873 GHz

GLENN J. WHITE^{1,2} AND ROBIN DEY³

Version October 4, 2026

ABSTRACT

We ask whether any nearby star operates a narrowband millimetre or submillimetre transmitter bright enough for ALMA to have recorded it incidentally. Technosignature searches have barely touched 90–873 GHz; published work lies almost entirely below 30 GHz.

We searched the ALMA archive for an unresolved, drifting carrier at a stellar position. The sample is 1651 windows toward 82 systems within 40 pc, all observed for other science, mostly debris discs. Of these, 402 windows toward 60 systems have channels fine enough to follow a drifting carrier. These cover 47.7 GHz of sky frequency, in disconnected intervals, not one continuous band.

Sensitivity was measured, not modelled, by injecting artificial carriers into the calibrated visibilities and recovering them through the same pipeline. The median system’s best window reaches $\text{EIRP}_{90} = 1.77 \times 10^{15}$ W. A 12 m dish radiating 1 MW at 230 GHz appears at 5.9×10^{14} W, reached for 14 of the 60 systems.

The search returns 56 threshold crossings, and the unexplained ones are as numerous as chance alone predicts. Carbon monoxide accounts for 19, including the β Pictoris disc, recovered with no prior information. Only 1 crossings outrank all 512 spatial controls, and neither recurs at a second epoch. Stacking epochs for 69 stars in the stellar rest frame reaches 1.73×10^{13} W for Proxima Centauri, for carriers whose reflex drift stays within 20.9 km s^{-1} .

We find no persistent narrowband transmitter. ALMA chose these targets, so the result constrains these 82 systems, not the solar neighbourhood.

Subject headings: extraterrestrial intelligence, radio lines: stars, submillimetre: stars, surveys, astrobiology

1. INTRODUCTION

In plain terms, this paper asks whether any nearby star is transmitting a narrow radio signal that ALMA happened to record while observing something else. ALMA’s public archive holds thousands of hours of millimetre and submillimetre observations taken for quite different reasons, mostly studies of dust discs and planet formation. Every one of them also contains a spectrum of the star at the centre of the field. A transmitter would appear in such a spectrum as a spike confined to one or two channels. It would sit at the star’s position and nowhere else in the field, and reappear if the star were observed again. We searched the archive for exactly that, with one pipeline and no human intervention. Nothing of the kind survived.

A technological civilisation may imprint the electromagnetic spectrum detectably whether or not it intends to, a *technosignature* (Sheikh 2020; Vidal et al. 2026). Millimetre wavelengths are an attractive place to look for one. Interstellar scattering and dispersion fall steeply with frequency, so a high-frequency carrier arrives less distorted. A given aperture delivers higher gain, and so more radiated power per watt supplied to the transmitter. Coherent millimetre emission is also already engineered on Earth, for directional links, radar and beamed

power.

The band is correspondingly hard to search. The atmosphere transmits unevenly and adds phase noise, and receivers are noisier than at centimetre wavelengths. Archived correlator channels are thousands of times coarser than a dedicated narrowband backend, so a carrier is diluted within one channel and its drift can smear it across several. These bands are also crowded with molecular lines, a dense background of real astrophysical signals that a threshold search must separate from an artificial one.

Six decades of radio searches have examined many thousands of stars, yet sample only 6×10^{-18} of the space of transmitter frequency, bandwidth, power and duty cycle (Wright et al. 2018). Essentially all of that effort lies below ~ 30 GHz; above it we are aware of one published search, in ALMA Band 3 (Mason et al. 2025). Few searches at any frequency reprocess existing interferometric archives, although those archives already hold a spectrum of every target they have pointed at (§2).

This paper is an archival ALMA search of nearby stellar systems across that band, with its sensitivity and its selection measured rather than assumed. The signal class is an *unresolved spectral carrier*: emission confined to a few channels, unresolved at the stellar position, and persistent enough to recur between epochs. Each qualifier is a real restriction. A carrier broader than the channel width would not be recovered here, nor one that fills the synthesised beam, nor one that appears once and never again. The fine channels used are 15.3–1953 kHz wide, so “narrowband” here describes the assumed transmitter

¹ School of Physical Sciences, The Open University, Walton Hall, Milton Keynes, MK7 6AA, UK; glenn.white@open.ac.uk

² RAL Space, STFC Rutherford Appleton Laboratory, Harwell Campus, Didcot, OX11 0QX, UK

³ VBRL Holdings, 50/5 Huay Kaew Road, Chang Phueak, Chiang Mai, Thailand

and never the resolving power of the data. Channelisation also divides the archive. Only 402 of the 1651 searched windows support the fine-channel search; the rest register an unexplained spike but cannot tell a drifting signal from a stationary one.

Every window follows the same four steps: a threshold crossing; attribution to a known line; comparison of the star with 512 control positions in the same field; and a search for recurrence at another epoch. Figure 2 is the roadmap. One worked case runs through the paper: a crossing toward β Pictoris, the largest star-versus-control contrast here, attributed to the disc’s known circumstellar CO in two transitions and 9 execution blocks (§5.3). Where this paper says “sensitivity” it means EIRP_{90} , the equivalent isotropic radiated power that the complete automated selection recovers nine times in ten. The internal trigger level P_{trig} is not a sensitivity.

Three limits on interpretation hold throughout. The targets were chosen by archival disc and planet-formation programmes, not for this search. The sample is therefore enriched in F and G stars and depleted in M dwarfs: only 21 of 5908 catalogued M dwarfs within 40 pc appear in it. Every statistical statement here concerns the ALMA archive-selected stellar sample, not the solar neighbourhood. The survey cannot constrain the local habitable-zone planet population, which M dwarfs dominate (§3.1). The search statistic and the line mask were defined on these data, so the rates quoted here are post hoc, and §4 sets out what was fixed when. And only 42 of 82 systems were observed more than a day apart, so a non-detection constrains a persistent carrier far better than an intermittent one.

2. BACKGROUND

Radio searches for technosignatures began at centimetre wavelengths and have largely stayed there. The first targeted experiments (Cocconi & Morrison 1959; Drake 1961), the programmes reviewed by Tarter (2001) and Project Phoenix (Backus & Project Phoenix Team 2002) all worked below a few gigahertz. The modern generation divides into targeted and commensal surveys at centimetre and metre wavelengths (Zhang et al. 2020; Czech et al. 2021; Morrison et al. 2023; Li et al. 2026; Johnson et al. 2023), with wide-field low-frequency coverage from the MWA (Tremblay et al. 2024) and single-dish C- and K-band work at 6 and 18 GHz with the Sardinia Radio Telescope (Manunza et al. 2025). Machine-learning interference rejection is now central to that effort (Ma et al. 2023; Choza et al. 2024; Jacobson-Bell et al. 2025; Poznanski et al. 2025; Garrett et al. 2026). The classical threshold search used here lacks that capability, and instead uses a spatial test on the star against control positions in the same field, which a single dish cannot make.

Above about 30 GHz the record is nearly empty, and that is what makes this frequency range new. ALMA observes in Bands 3 to 10, spanning roughly 84 to 950 GHz; of those eight bands, one has been searched for technosignatures. For continuity with the centimetre literature we note one familiar power scale: the Arecibo planetary radar radiated 2.0×10^{13} W at 2.38 GHz (Drake 1974; Sheikh et al. 2025). It carries no transmitter model into the millimetre band, where an aperture of that diameter is not realisable, so the illustrative benchmark of §6.3

is instead a dish of 12 m. Neither is a privileged scale; the limits of this paper are stated as EIRP and need no transmitter model.

Mason et al. (2025) is the only prior millimetre search and the only comparable archival one. They searched Band 3 windows toward 28 Gaia DR3 stars lying in the fields of 4 ALMA calibrators, observed for other science and so usable for free. Their nearest target is at 1.01 kpc, where a 5σ threshold corresponds to $\text{EIRP}_{\text{min}} > 6.9 \times 10^{17}$ W. Their targets are four orders of magnitude more distant than ours, so an EIRP ratio between the two surveys would be almost entirely a distance term. The quantity that does not depend on distance is the minimum detectable *received* flux, $F_{\text{min}} = S_{\text{min}} \delta\nu$. It must be quoted at a matched channel width, because S_{min} falls as $\delta\nu^{-1/2}$ while F_{min} rises as $\delta\nu^{1/2}$. Mason et al. reach 3.8×10^{-23} – 5.8×10^{-23} W m $^{-2}$ across their four fields, on a 30.52 kHz channel at a bare 5σ threshold. On that same channel width and the same threshold convention, the fine-channel windows searched here reach 5.2×10^{-24} W m $^{-2}$ at best, a factor 9 deeper, and 7.1×10^{-23} W m $^{-2}$ at the median, a factor 1.5 shallower. The two experiments are therefore comparable in depth and differ in what they cover: distant stars in four calibrator fields, against 60 stellar systems within 40 pc searched for a drifting carrier, over a frequency range reaching about nine times higher. They also identified the obstacles that shape the present method: drift rates scale with frequency, smearing erodes sensitivity unless drift is searched for, and the dense forest of molecular lines raises the confusion background. Every window reported here was re-reduced and re-searched from the raw archive data with the pipeline of §4, so for any star in common the numbers to cite are this paper’s.

Figure 1(a) places this survey among published searches, in minimum detectable EIRP against distance. The vertical axis is physically commensurable: for a carrier narrower than every programme’s channel, the minimum detectable EIRP is a total radiated power and not a spectral power density, so the same quantity is plotted for each programme. What differs is the convention. The value plotted for this survey is EIRP_{90} , a 90 per cent recovery through the complete automated selection, whereas most published limits are nominal thresholds at which recovery is not stated; the comparison runs against us. Any EIRP advantage toward the nearest systems is a d^2 effect and not an instrumental one. Panel (b) shows the other axis that matters. Archived correlator channels here span about seven orders of magnitude in width, and all of them are far coarser than a dedicated narrowband backend. The distinctive feature of this survey is therefore its frequency regime and its channelisation, not its depth. It also samples far fewer systems and far fewer epochs than a dedicated search, and complements those programmes rather than deepening them.

A non-detection becomes a quantitative statement only through a stated search volume. Wright et al. (2018) formalised that volume, and Margot et al. (2023) its conversion into an upper bound on the fraction of stars hosting a transmitter; Sheikh et al. (2025) asked where Earth’s own technosignatures would be detectable. Such a bound is valid only if the sample represents the population it is meant to describe. This sample was assembled by other




FIG. 1.— Where this survey sits among published radio technosignature searches. (a) Minimum detectable EIRP against distance: blue, this survey; the deepest window per star; orange, Enriquez et al. (2017); aqua, Margot et al. (2023); violet, Mason et al. (2025); left-edge ticks, Price et al. (2020). Grey dotted, $\text{EIRP} \propto d^2$ at constant received flux; green dashed, the illustrative 12 m benchmark transmitter of §6.3. The ordinates are commensurable as total powers for a sub-channel carrier; what differs between programmes is the detection threshold and the completeness level each adopts, and this survey’s value is the stricter of the two conventions. (b) Native channel width against observing frequency for every searched window (filled, Class A; open, Class B), with the comparison programmes as horizontal segments. The channel widths span about seven orders of magnitude, which is what the $\Delta \log \Delta \nu_{\text{ch}} \approx 7$ annotation marks: the channelisation, not the plotted power, is what differs most between programmes.

observers for other purposes, so it does not, and §3 states what it is instead.

3. THE SAMPLE

This search covers 1651 spectral windows toward 89 stars in 82 systems within 40 pc, drawn from 404 public ALMA execution blocks and 4.1 TB of raw data through

one frozen pipeline, h on source. We call that set the *primary census*, and it is the set every limit, every crossing and every expectation in this paper is computed over. Four further sets of blocks enter the paper — a hold-out reserved before the archive was searched, an archival extension searched afterwards, an out-of-sample control beyond 40 pc, and an external sample used only to calibrate the spatial statistic — and none of them is part of the primary census. Table 2 gives all of them in rows that sum, with the results each one enters; wherever a number in this paper is computed over something other than the primary census, that table’s name for it is used at the point of use.

Three words are used consistently and are not interchangeable. A *star* is an individual stellar component; a *system* is a gravitationally bound group, counted once however many of its components were observed; a *target* is one ALMA pointing. The census holds 89 stars in 82 systems.

The census is two experiments, not one. 402 windows toward stars in 60 systems are fine-channel, called *Class A* throughout: these resolve a drifting carrier and are the primary experiment, and every sensitivity and every exclusion quoted here refers to them unless stated otherwise. The remaining 1249 windows are coarse-channel, *Class B*, and support a secondary search for unresolved spectral excess; they can register a spike but cannot distinguish a drifting signal from a stationary one.

The inclusion criteria are objective and were applied by query. A star enters if its Gaia DR3 parallax is at least 25 mas, measured to better than 20 per cent, and if its proper-motion-corrected position at the observation epoch falls inside the quoted field of view of at least one public execution block. Distances are direct inversions, $d = 1/\varpi$, and the star need not sit at the pointing centre. Disc, planet and spectral type enter as descriptive metadata only, and no star or window was excluded by an unlisted criterion or for short integration. The parallax criterion never binds: the worst fractional parallax error among the stars searched is per cent, an order of magnitude inside it. The query returns 168 work-list entries; the quoted field of view is optimistic, and only 115 of those stars are in fact contained by a public field, of which 89 yielded a searchable window (Fig. 2). The per-star and per-window ledger, including the five-criterion decision tree and every edge case, is in the data release.

3.1. The archival selection

This sample was not designed. It is whatever previous ALMA programmes happened to observe, for their own reasons, and 90 per cent of the stars searched were targeted under a debris-disc or planet-formation proposal. The resulting bias is strong and measurable, and Table 1 tabulates it against a volume-limited reference census selected on the same parallax quality.

The dominant selection effect is distance. Stars within 5 pc are over-represented by a factor , and the enhancement falls monotonically with distance to a factor of two under-representation beyond 30 pc. Spectral type comes second: A stars are over-represented by about and F and G stars by factors of two to three, because those are the stars with bright debris discs. M dwarfs are mildly depleted — of the searched stars are M dwarfs, a factor of their share of the reference census, and only of them

TABLE 1

SELECTION FUNCTION OF THE SEARCHED SAMPLE AGAINST A VOLUME-LIMITED REFERENCE CENSUS, BOTH SELECTED ON THE SAME PARALLAX QUALITY: GAIA DR3 SOURCES WITH $\varpi > 0$, $\sigma_{\varpi}/\varpi \leq 0.1$ AND $1/\varpi \leq 40$ PC (17 566 OBJECTS, 8594 OF THEM CARRYING A `TEFF_GSPPHOT` ESTIMATE), AGAINST THE 89 STARS WITH AT LEAST ONE SEARCHED WINDOW. THE SEARCHED STARS ARE MEASURED FAR BETTER THAN THE CUT REQUIRES — THE WORST FRACTIONAL PARALLAX ERROR AMONG THEM IS 2.5 PER CENT — SO THE TWO COLUMNS ARE SELECTED ALIKE, AND LOOSENING THE CUT TO 0.2 ADDS NO CENSUS OBJECT AT ALL. 86 OF THE 89 SEARCHED STARS ARE CLASSIFIED: 53 FROM THE GAIA TEMPERATURE, 25 FROM THE SIMBAD SPECTRAL TYPE WHERE THAT TEMPERATURE IS ABSENT, AND 8 IDENTIFIED AS WHITE DWARFS; 3 CARRY NO CLASSIFICATION IN EITHER SOURCE.

THE SIMBAD ROUTE MATTERS BECAUSE `TEFF_GSPPHOT` IS MISSING PREFERENTIALLY FOR COOL STARS, WHICH IS THE POPULATION THE M DWARF STATEMENT IS ABOUT. “RATIO” IS THE SAMPLE COLUMN FRACTION OVER THE CENSUS COLUMN FRACTION; > 1 IS OVER-REPRESENTATION. *The strongest selection effect here is distance, not spectral type.*

Property	Census	Searched	Ratio
Distance			
$d \leq 5$ pc	48 (0.3%)	12 (13.5%)	49.3
5–10 pc	240 (1.4%)	13 (14.6%)	10.7
10–20 pc	2069 (11.8%)	18 (20.2%)	1.7
20–30 pc	5391 (30.7%)	19 (21.3%)	0.7
30–40 pc	9818 (55.9%)	27 (30.3%)	0.5
T_{eff} class ^a			
A	24 (0.3%)	5 (6.4%)	23.0
F	402 (4.7%)	9 (11.5%)	2.5
G	860 (10.0%)	16 (20.5%)	2.0
K	1400 (16.3%)	7 (9.0%)	0.6
M	5908 (68.7%)	41 (52.6%)	0.8
Other axes			
Confirmed planet hosts ^b	26 (15.5%)	23 (25.8%)	1.7
White dwarfs ^a	—	8 (9.0%)	—
Searched stars / systems	—	89 / 82	—

^aTemperature bins, not MK types: A ≥ 7500 , F 6000–7500, G 5300–6000, K 3900–5300, M < 3900 K, normalised over the 78 classified non-degenerate stars. The 8 white dwarfs are kept apart because a degenerate is not an early-type star and its photometric temperature would place it among them. Gaia `teff_gspphot` is unavailable for about half the reference census, preferentially for the coolest and faintest objects, so the census M fraction is a lower limit and every earlier-type ratio in this panel is correspondingly a lower limit; the searched column has no such gap.

^bBoth columns are the same positional join, of the 168-entry ALMA-covered work list and of its searched subset, to the NASA Exoplanet Archive `pscomppars` table accessed 2026 October 4; a confirmed planet is a row of that table. The 23 searched hosts carry 50 confirmed planets between them, and 21 of the 23 survive the deuterium-burning mass limit. The Gaia reference census carries no planet information, so the census column in this row is the work list rather than the 17 566-object Gaia census.

carry the fine-channel coverage that can resolve a drifting carrier at all — and that is a weaker statement than the table’s early-type ratios, not a stronger one. Both figures are bounds in the same direction: Gaia temperatures are missing for about half the reference census and missing preferentially for the coolest objects, so the census M fraction is a lower limit and every early-type ratio is therefore a lower limit too. Coverage is also concentrated: 285 of the 1651 windows come from 3 systems. Every statistical statement in this paper is therefore a statement about an archive-selected sample of nearby stars, and none of them is an occurrence rate for the solar neighbourhood. The systems a targeted search would itself have chosen are collected in Table 3.

3.2. Frequency coverage and archival scope

Selection




FIG. 2.— The selection funnel, from the catalogued population within 40 pc to the stars with a searched window. The three quantities the text keeps apart are the 168-entry work list the archive query returns, the 115 stars a public field genuinely contains, and the 89 that yielded a searchable window. The chain applied to what comes out of the funnel is Fig. 8.

The Class A experiment covers $\Delta\nu_A = 47.7$ GHz of unique sky frequency, and that coverage is sparse: it is 26 disjoint islands scattered between 114.0 and 873.1 GHz, not a band. The endpoints are a range over which observations exist, never a range that was covered. Adding the coarse-channel windows brings the total archival coverage to 118.1 GHz in 46 intervals, of which 113.3 GHz survives the molecular-line mask of §4. The distribution is shown in Fig. 3; a transmitter at a frequency between the islands was never observable here, whatever its power.

The archive is not exhausted either, and the shortfall is accounted for rather than estimated. further blocks lie inside the same query but were not processed when the census was frozen; Table 2 gives what became of each. of them have since been calibrated and searched, and they enter this paper only as repeat epochs. The largest single gap is the 18 blocks for which the archive holds raw visibilities that were never pipeline-calibrated, 237 GB in all; no effort on our side would close it.

4. DATA AND METHODOLOGY

The search looks for one kind of signal: an unresolved carrier at the star’s own position, narrow compared with a single correlator channel, and drifting in sky frequency because the transmitter accelerates along the line of sight. Over a long track its power is spread across many channels, so a search that treats each channel on its own will miss it. The analysis de-drifts the data before it measures them.

4.1. The de-drifted search statistic

Each window’s dynamic spectrum is stacked along every trial drift rate $\dot{\nu}$. The search statistic is the largest stacked signal-to-noise ratio anywhere on the resulting channel $\times$ drift grid,

$$T(\mathbf{x}) = \max_{\dot{\nu}, \nu} \frac{\sum_t I(\mathbf{x}, t, \nu + \dot{\nu}t) / \sigma^2(t, \nu)}{[\sum_t 1 / \sigma^2(t, \nu)]^{1/2}}, \quad (1)$$

where $I(\mathbf{x}, t, \nu)$ is the spectrum at sky position $\mathbf{x}$ in integration t and channel ν , and $\sigma(t, \nu)$ is the noise scale




FIG. 3.— Where this survey searched. *Lower panel*: each searched window is a horizontal segment on its host system's row, rows ordered by distance and labelled alternately on the left and right axes; heavy blue marks the fine-channel carrier search (Class A) and light grey the coarse-channel spectral-excess search (Class B). *Upper panel*: the number of systems covering each gigahertz of sky frequency; the threshold distribution is Fig. 1(a). The coverage is narrow, clustered and repeatedly revisited: 46 islands concentrated near 230 and 345 GHz, with summed window bandwidth 23.9× the union. The union measures frequency coverage, the summed bandwidth integrated search effort.

there. Equation 1 is the matched filter for a constant-drift carrier: it adds the signal coherently along the track and divides by the noise of the stacked spectrum, so T is an absolute signal-to-noise ratio and not a local contrast. The noise scale is the median absolute deviation across 512 control positions at each integration and each channel — local in time and frequency, global in position, and shared by the star and every control, with the across-position median never subtracted from the numerator. The statistic is evaluated at one position at a time, identically at the star and at each control, which is what makes the two comparable. T_* denotes its value at the star.

Spectra are extracted at the star's apparent position,

never from images. For each integration and each native channel the pipeline forms the weighted vector average of the phase-shifted visibilities, in total intensity: the mean $I = (XX + YY)/2$ of the two parallel-hand correlations. That average is the dirty-map value at the extraction position: an amplitude, not a deconvolved flux density, so in a field with extended emission it is not a point-source measurement. The data are used as the archive delivered them, through ALMA's own calibration and quality assurance, with their flagging carried through and with no regridding and no deconvolution. All delivered channels are searched, with no edge-channel trimming. A per-integration median along the *frequency* axis removes the continuum; no median is taken along time, so a one-

TABLE 2

THE BLOCK AND WINDOW LEDGER. EVERY SET OF EXECUTION BLOCKS USED ANYWHERE IN THIS PAPER, WITH THE RESULTS IT ENTERS; THE ROWS SUM, AND EACH INDENTED GROUP SUMS TO THE ROW ABOVE IT. THE PRIMARY CENSUS IS THE SURVEY: UNLESS A STATEMENT NAMES ONE OF THE OTHER SETS, IT IS COMPUTED OVER THAT ROW. HOURS ARE ON-SOURCE HOURS. CONFIRMED PLANETS IN TABLE 3 ARE TAKEN FROM THE TABLE, ACCESSED .

Set of execution blocks	Blocks	Windows	Stars	Systems	Hours	Results it enters
Archival scope: blocks the query exposes	656	—	—	—	—	defines the scope; enters no statistic
Processed with the frozen pipeline	484	—	—	—	—	the sum of the three rows below
Primary census	404	1651	89	82	1094	every limit, the crossing list, every chance expectation
Class A, fine channels ^a	—	402	65	60	227	the drifting-carrier experiment
Class B, coarse channels ^a	—	1249	88	81	867	the spectral-excess experiment
Reserved hold-out	77	315	36	—	214	the rank displacement, the false-alarm tail factor, the radial control profile
No window survived the quality cut	3	0	0	0	0	nothing
In scope, not processed at the freeze	177	—	—	—	—	the sum of the five fates below
Calibrated and searched	144	—	—	—	—	recurrence only; never a denominator
Archival extension , reported	115	463	16	16	269	recurrence, and the null on data the frozen pipeline had not seen
No calibration in the ALMA delivery	18	—	—	—	—	nothing; the archival gap
Calibration failed	11	—	—	—	—	nothing; the archival gap
Not attempted	3	—	—	—	—	nothing; the archival gap
Working disk insufficient	1	—	—	—	—	nothing; the archival gap
Beyond 40 pc: out-of-sample control	47	198	22	—	88	the external null; no science count
External calibration sample ^b	326	1322	56	—	—	the rank calibration and the positive control; no science count

^aThe two classes partition the census windows, not its blocks: 252 blocks carry a Class A window and 389 a Class B one, and most carry both.

^bNot a subset of the archival scope: it includes blocks toward stars outside the 40 pc work list, so its blocks cannot be placed inside this ledger.

TABLE 3

SYSTEMS OF PARTICULAR INTEREST: EVERY SAMPLE SYSTEM WITHIN 15 pc HOSTING AT LEAST ONE CONFIRMED PLANET, ORDERED BY DISTANCE. N_{pl} IS THE CONFIRMED-PLANET COUNT, N_{A} THE NUMBER OF CLASS A WINDOWS, EIRP₉₀ THE DEEPEST CLASS A SENSITIVITY REACHED, AND N_{eb} THE NUMBER OF PRIMARY-CENSUS EXECUTION BLOCKS COVERING THE SYSTEM, SO $N_{\text{eb}} > 1$ MEANS REPEAT EPOCHS EXIST WITHIN THE CENSUS. SIX OF THESE SYSTEMS HAVE NO CLASS A COVERAGE AND SO CARRY NO DRIFT-RESOLVING LIMIT — THE ARCHIVAL SELECTION OF §3.1 ACTING ON EXACTLY THE TARGETS A TRANSMITTER SEARCH WOULD CHOOSE. THE FULL TARGET LIST IS IN THE DATA RELEASE.

System	d (pc)	N_{pl}	Bands	N_{A}	EIRP ₉₀ (W)	N_{eb}
Proxima Cen	1.3	2	6	0	—	25
Barnard’s Star	1.8	4	6	0	—	3
τ Cet	3.7	3	6	6	1.7×10^{14}	8
BD+05 1668 (GJ 273)	3.8	2	6	0	—	25
HD 33793	3.9	1	6	0	—	22
GJ 674	4.6	1	6	0	—	20
HN Lib	6.3	1	6	5	5.6×10^{14}	5
GJ 581	6.3	3	6	0	—	9
61 Vir	8.5	3	7	5	1.8×10^{14}	5
GJ 849	8.8	2	6	6	9.9×10^{14}	6
HD 285968	9.5	1	6	4	1.3×10^{15}	4
AU Mic	9.7	4	3,6,9	13	1.7×10^{14}	14
TRAPPIST-1	12.5	7	3,6,7	5	2.4×10^{14}	18
HD 69830	12.6	3	6,7	22	1.7×10^{14}	8
LHS 1140	15.0	2	6	4	2.8×10^{15}	4

or two-channel feature cannot be removed by it. One subtlety matters for every power quoted below. The delivered channels are smoothed, so a channel’s effective noise bandwidth is not its separation, and flux density times channel width is not yet a radiated power. A response correction computed from the correlator’s own channel shape supplies the difference. It is derived, and validated against the effective resolutions ALMA reports, in Appendix A.

One convention is fixed here and holds everywhere in the paper. The search itself runs on the topocentric sky frequencies ALMA delivered, and every drift rate quoted is the apparent one at the telescope, since that is what

TABLE 4

WHAT THE SURVEY SEARCHED WITH. COUNTS OF BLOCKS, WINDOWS, STARS AND SYSTEMS ARE IN TABLE 2 AND ARE NOT REPEATED HERE. EVERY FREQUENCY UNION QUOTED IS AN AGGREGATE OVER CLUSTERED, REPEATEDLY REVISITED COVERAGE, WHOSE DISTRIBUTION IS FIG. 3; NO UNION SHOULD BE READ AS CONTINUOUS COVERAGE BETWEEN ITS ENDPOINTS.

Quantity	Value
Distance range	1.30–39.63 pc
Windows on the 12 m array; on the ACA	597; 1054
Native channel widths	15.3 kHz–31.25 MHz; median 15.625 MHz
$\Delta\nu_{\text{A}}$, unique Class A frequency union (the drifting-carrier experiment)	47.7 GHz in 26 islands, 114.0–873.1 GHz
coverage Class B adds	70.4 GHz
total archival coverage	118.1 GHz (46 intervals)
usable after the molecular-line mask	113.3 GHz
Nominal trigger P_{trig} , Class A;	2.2×10^{13} – 6.9×10^{16} W;
Class B	1.6×10^{13} – 2.8×10^{17} W

the data contain and what a reader can recompute from the released channel axes. Two steps need the star’s own frame rather than the sky’s, and both take it: a crossing is compared with the line catalogue, and a repeat epoch is read at the predicted cell, after each block’s frequency axis has been corrected for the Earth’s motion at that block’s epoch and then for the star’s systemic velocity. Nothing in this paper is evaluated in a barycentric frame alone.

Two power quantities carry the main text. The trigger level P_{trig} is the equivalent isotropic radiated power, $\text{EIRP} = 4\pi d^2 S \Delta\nu$ at distance d , of a carrier that makes $T_{\star}$ reach $T_{\text{trig}} = 5$ in one native channel of width $\Delta\nu$. It is an internal pipeline level, the power at which the search fires, and it is neither a sensitivity nor a significance. The sensitivity is EIRP₉₀, the power at which nine injected carriers in ten are recovered through the complete selection of §4.2, the spatial ranking requirement included. It is a recovery fraction through this pipeline, measured rather than assumed, and it is not

the probability of a statistically significant detection: no stage of the chain carries a calibrated significance, so no such probability exists to quote. 8860 synthetic carriers of known amplitude were injected into calibrated visibilities and recovered through the unmodified pipeline. They span 52 block-window configurations, both ALMA arrays, four bands and the sample's full range of channel widths. For Class A the result is $\text{EIRP}_{90} = 4.42 P_{\text{trig}}$. Every limit quoted in this paper is an EIRP_{90} , given to two significant figures beside its native channel width; its error budget is Table 5. The absolute flux scale has been checked twice from outside. Our extractor reproduces an independent reduction of the same AU Mic visibilities to 0.994 ± 0.014 on a 29.4 mJy point source. It reproduces the published quiescent flux density of ϵ Eridani to 1.10 ± 0.31 (Burton et al. 2022). The two checks share neither band nor array.

The drift range searched is bounded by an acceleration, not by a frequency. Drift rate and line-of-sight acceleration are one quantity in two sets of units, $a_{\text{los}} = c \dot{\nu}/\nu$, so a ceiling expressed as $\dot{\nu}/\nu$ is the same physical limit in every band. The default ceiling is $12.0\text{--}13.3 \text{ Hz s}^{-1} \text{ GHz}^{-1}$, that is $|a_{\text{los}}| \leq 3.60\text{--}4.00 \text{ m s}^{-2}$. That is the acceleration of a transmitter orbiting a solar-type star at a few hundredths of an au. It is the Mason et al. (2025) fiducial with a threefold margin, widened for a confirmed exoplanet host whenever the innermost known planet demands more. The margin is not cosmetic: at unit margin the search would have been blind to two of the windows that produced a crossing. The grid is uniform in $\dot{\nu}$, and fine enough that a carrier midway between two trials loses no more than 0.26 of a native channel over the longest track. Two limits are worth stating plainly. The rates searched are the apparent ones at the telescope, so a deliberately chirped carrier is not distinguished from an accelerating platform. And a transmitter need not be bound to a known planet at all, so drifts faster than the ceiling are simply unconstrained (Fig. 4).

Only part of the archive can carry a drift search at all, and the two classes reported here are that division. Channelisation in the sample is bimodal, with an empty gap between the modes, so the split is unambiguous. Class A is the primary experiment: the 402 fine-channel windows, median 488.3 kHz, in which the maximum searched drift crosses of order one channel or more over the on-source track ($\eta_{\text{drift}} = |\dot{\nu}_{\text{max}}| \tau_{\text{track}} / \Delta \nu_{\text{ch}} \gtrsim 1$), and the only class with measured end-to-end recovery of a drifting carrier. Class B is the secondary experiment: the 1249 coarse-channel windows, mode 15.625 MHz, in which a carrier stays inside one channel for the whole track; these constrain a persistent unresolved spectral excess, carry no drift discrimination, and are never described as a carrier search here. Each window's own η_{drift} , trial-drift count and acceleration ceiling are released catalogue columns, so a reader may reclassify on the physical criterion instead; the two criteria do not partition the survey identically. Results are given per class, and combined only where that is statistically appropriate.

4.2. From a crossing to a candidate

Every window passes through the same four tests in the same order, and three words name what comes out of them: a *crossing*, a *candidate* and, had one survived,




FIG. 4.— What the drift search reaches, as a physical selection function. The drift ceiling of each of the 402 Class A windows (points), converted to line-of-sight acceleration, $|a| = c \dot{\nu}/\nu$, against observing frequency; the shaded region above is unsearched. The boundary is flat because the grid was built at constant acceleration; in the units the search works in it is not, running 1.4 kHz s^{-1} at 115 GHz to 10.5 kHz s^{-1} at 872 GHz. Horizontal lines are reference accelerations: Earth's rotation (0.034 m s^{-2} at the equator) and orbit (0.0059 m s^{-2}), and the two most demanding known planets of any target, TRAPPIST-1 b (4.00 m s^{-2}) and GJ 581 e (2.27 m s^{-2}). The boundary drawn is the default ceiling, 3.60 m s^{-2} , which is also the median; it runs to 4.00 m s^{-2} on the 5 TRAPPIST-1 windows, and it is each window's own ceiling that decides whether a given orbit is contained. The most demanding case is TRAPPIST-1 b, whose line-of-sight acceleration exceeds the default ceiling; TRAPPIST-1 transits, so the edge-on figure is the relevant one, and planet b lies above the default ceiling for 29 per cent of its orbital phase and above the widened one for 1 per cent. It is searchable through the rest of the orbit, and planet c reaches the default ceiling only at modest eccentricity.

a *confirmed* signal (Fig. 8). Each test is set out below as the quantity it measures, the number that quantity is compared with, when that number was fixed, and what passing it does and does not establish. There is no fifth test, and not one of the four is a significance test.

- (i) *Trigger* $\rightarrow$ a *crossing*. The quantity is T_* , the drifted statistic of Eq. 1 read at the stellar position; the criterion is $T_* \geq T_{\text{trig}} = 5$ in at least one channel $\times$ drift cell; and both the statistic and the level were fixed before the archival search began. **T_{trig} is a trigger threshold and not a detection significance.** T is not shown to follow any standard distribution under the null, so the level carries no tail probability, and a window holds between 112 and 6.87×10^6 independent cells, so one level buys very different protection in a coarse single-pointing window and in a finely channelised one on a long track. A crossing establishes only that the window goes on to the next three tests.
- (ii) *Line attribution*. The quantity is the crossing's velocity offset from the nearest transition of a frozen molecular-line catalogue, evaluated in the star's own rest frame; the criterion is $\pm 50 \text{ km s}^{-1}$, inside which the crossing is astrophysical and goes no further; and the catalogue, the species list and the half-width were all fixed before any crossing was classified. The half-width follows the Keplerian line-of-sight velocities of gas surviving beyond a few au, $\lesssim 13 \text{ km s}^{-1}$, with margin for the systemic velocity. The mask classifies rather than detects: it is applied after the search, so no threshold moves and no window goes

unsearched. What it establishes is a kinematic coincidence with a known line, and its price is that the survey constrains nothing inside the excluded intervals — ≈ 113.3 GHz of the 118.1 GHz covered survives the mask (Appendix C).

- (iii) *Spatial selection* $\rightarrow$ *a candidate*. The quantity is T_* ranked against the same statistic at 512 control positions in the same window, drawn once from a fixed seed into an annulus placed by each block’s own synthesised and primary beams so that no control sits on the star. The rank runs the way a competition does: its *lowest* value, $1/(512+1)$, means the star beat every control, so a rank below one half places the star high in its own field. The criterion is that lowest value, and the ring, its seed and the criterion were all fixed before any crossing was inspected. **This is a localisation and ranking filter, not a significance test.** The stellar and control positions are not demonstrably exchangeable, so no per-window false-alarm probability can be attached to the rank; and the filter responds to compactness rather than to origin. The clearest demonstration is β Pictoris, where genuine carbon monoxide fails the screen because the resolved disc lifts the control positions along with the star. *The screen measures compactness rather than authenticity*, which is why it decides which crossings are examined and never establishes one. Requiring it costs a factor $\times 1.44$ in sensitivity, from $3.06 P_{\text{trig}}$ at the trigger alone to $4.42 P_{\text{trig}}$ through the screen, and every completeness in this paper is charged for it (Appendix B).
- (iv) *Independent recurrence* $\rightarrow$ *confirmation*. The quantity is T_* at the cell the first epoch predicts, read in a later execution block of the same star at a single drift rate after the frame correction of §4.1; the criterion is that it reach the trigger again; and the test was specified before the repeat epochs were searched. **This is the survey’s only confirmation criterion**, and the archive’s uneven temporal sampling is the main reason it is a weak one: only part of the sample has epochs separated well enough to apply it at all.

Most crossings are expected by chance, because a fixed threshold is not a fixed false-alarm rate. The expectation is computed window by window, from that window’s own cell count and the exceedance rate measured at its own control positions, then summed over the windows searched. It predicts 34.5 unattributed Class A crossings against 35 observed, about two-thirds of the 52 Class A crossings. Appendix D gives the conditioning of that expectation and the control ensemble it is measured on.

The displacement that leaves the rank uncalibrated is measured out of sample and independently of the crossings: on 315 windows the frozen pipeline had never seen, the stellar rank has median 0.405 where one half is expected if the star and its controls were interchangeable, and a parameter-free correction for the annulus geometry does not restore an exchangeable statistic (Appendix B). Throughout this paper p denotes a genuine significance level, and no probability computed against the screen’s own null is called one.

Three corrections precede the statistic: the extraction field is chosen so that the star lies inside the delivered

primary-beam field rather than merely inside the quoted field of view; the stellar position is propagated to the observation epoch and used as the apparent position; and on-source time is accounted per integration from the delivered exposure records rather than from the scan intentions.

A visibility-domain check asks whether a crossing’s emission is at the stellar coordinates rather than elsewhere in the field. It is applied to every crossing but reported beside the chain rather than inside it, because it cannot discriminate at the trigger (§5.5.1).

That the chain finds what it is meant to find is shown twice over: by the injected carriers above, and by β Pictoris CO, a real feature at a known stellar position, recovered and localised in two transitions and 9 execution blocks (§5.4).

5. RESULTS

The primary experiment is the Class A search for a drifting, unresolved spectral carrier: 402 fine-channel windows toward 60 systems. The Class B search for an unresolved spectral excess is secondary and carries no drift discrimination: 1249 coarse windows. Together they are the 1651 windows of the released catalogue. **No carrier was found toward any of the 82 systems searched.** What follows is first how deep the search reaches, and then what it found and how each finding was disposed of.

5.1. Sensitivity: the power at which a carrier would have been found

The sensitivity of this survey is a measured completeness and not a threshold. Artificial carriers were added to the real calibrated visibilities, each at a random sub-channel phase and with a random drift rate from its window’s own grid, and recovered through the unchanged pipeline: extraction, baseline removal, de-drifted stack, trigger and the 512-position spatial screen. The power at which nine carriers in ten come back is EIRP_{90} , the only sensitivity quoted here. It charges for the whole chain, because a carrier counts as recovered only when it reaches the trigger and outranks every control.

Measured that way the completeness is not one number, and what it depends on is the control ring. In the 390 of 402 Class A windows whose ring is noise, nine carriers in ten are recovered at $3.45 P_{\text{trig}}$, where P_{trig} is that window’s own 5σ trigger power. In the other 12, toward β Pictoris, ALMA J1537–3319, HD 48370 and HD 139084, the ring carries resolved disc emission; a carrier there has to outshine the star’s own surroundings rather than the noise, and the injected ladder never reaches ninety per cent below $8 P_{\text{trig}}$, although every carrier is recovered by $20 P_{\text{trig}}$. The second figure is therefore a bound and not a value, and each window is quoted on the stratum it belongs to. A single factor fitted to both strata together would be 4.42, and that is the completeness of no window in the survey: it sits above the first stratum’s own spread and far below the second’s bound. Class B keeps one factor, $4.70 P_{\text{trig}}$: no injected coarse window has a bright ring, so the distinction cannot be measured there and the 4 coarse windows that do are flagged in the catalogue instead.

Separating the two strata moves power between windows; it does not make the survey better. The median

Class A window reaches 3.35×10^{15} W of equivalent isotropic radiated power and the median system, on its best window, 1.77×10^{15} W, with a spread across systems of 7.74×10^{13} – 5.79×10^{16} W (Fig. 6a). The survey reaches the 12 m benchmark transmitter of §6 (5.9×10^{14} W) for 14 of the 60 systems and the round 10^{15} W level for 22. Against a single blanket factor those limits are about a fifth deeper for 390 windows and several times shallower for 12, and the worst-placed system in the survey moves the wrong way because its best window is one of them. A second term runs against us throughout: the assumed phase centre omits the annual parallax, so a nearby star observed in a long-baseline configuration loses amplitude. The loss is 0.13 per cent in the median window and 51.5 per cent at worst, passing a tenth in 76 of the 1650 windows that carry the astrometric keys, and dividing each window by the fraction it retains makes every limit shallower — the median Class A limit by 1.5 per cent. One window, Wolf 28 in A002_X9f2ff8_X25b0, has no such keys and is left uncorrected. Table 5 collects these signed terms and the two-sided ones, which combine to $\times 0.77$ – $\times 1.12$ on every limit.

Carriers were injected directly into 21 Class A windows and the result transferred to the remaining 381, which is the dominant caveat on all of the above. Figure 5(a) shows the individual measurements rather than their summary: 18 of the 21 resolve a ninety per cent point, between 2.9 and $7.4 P_{\text{trig}}$, and the three that do not are the brightest-ring windows of all. There is no dependence on observing band or on channel width, the median values differing by 0.2 and 0.3. There is an apparent difference between the two arrays, 3.3 against 4.5, and it is not the array: it is carried by 4 windows whose target has bright emission near the ring, and it falls to 3.3 against 3.4 once the star’s own leak into the control positions is removed. That is why the strata are cut on the ring. Removing the leak is a diagnostic and not the screen we apply, since it would make the test for a source depend on a model of that source; and it buys little, moving the noise-dominated figure to 3.38 and the disc-affected windows not at all, because there the ring is set by the disc. Only the distance is worse for a single star than for the sample: of the 77 Gaia parallaxes the worst fractional error is 1.65 per cent, toward γ Tri.

5.2. The nearest stars: Barnard’s Star and Wolf 359

Barnard’s Star is the second-nearest stellar system, and Wolf 359 the third if brown dwarfs are excluded. To our knowledge these are the first millimetre or sub-millimetre technosignature limits for either; the searches checked are Enriquez et al. (2017), Price et al. (2020) and Margot et al. (2023). Both are non-detections. *Both limits come from the secondary Class B experiment and not from the carrier search:* the archive holds only coarse Band 6 windows toward these two stars, so neither is searched for a drifting carrier and neither constrains one. What is constrained is an unresolved spectral excess. Quoted on the EIRP₉₀ convention of §5.1, Barnard’s Star reaches 3.0×10^{14} – 4.5×10^{14} W over 12 windows. Wolf 359 reaches 3.5×10^{14} – 5.1×10^{14} W over 12.

5.3. Threshold crossings and how each was disposed of

The result is a chain of four tests, and it is read in one direction. Of the 1651 windows, 56 contain a chan-

TABLE 5
WHAT SETS EIRP₉₀, AND IN WHICH DIRECTION. THE TWO-SIDED TERMS COMBINE IN QUADRATURE TO ONE INTERVAL APPLYING TO EVERY LIMIT; THE SIGNED TERMS ARE BIASES, STATED WITH THEIR DIRECTION RATHER THAN FOLDED INTO AN ERROR BAR. THE CHANNEL RESPONSE IS CARRIED BY THE INJECTIONS RATHER THAN APPLIED AFTERWARDS: AT THE NINETY PER CENT POINT THE RECOVERED PEAK CHANNEL HOLDS $1.93\times$ THE NOMINAL TRIGGER, SO THE RESPONSE IS ALREADY INSIDE THE MEASURED COMPLETENESS AND APPLYING IT AGAIN WOULD COUNT IT TWICE.

	value	what it does to the limit
Measured completeness, as a multiple of a window’s own trigger power		
Ring is noise (390 windows)	$\times 3.45$	measured; $\times 3.37$ – $\times 4.20$ across resamples
Ring carries emission (12)	$> \times 8$	a bound: everything is recovered by $\times 20$, nothing at the trigger
One factor for both	$\times 4.42$	the completeness of no window; not used
Two-sided uncertainties (per cent), combined in quadrature		
Window-to-window transfer	$-22/ + 8$	the spread across the 21 injected windows
Flux, distance and visibility scale	± 9	calibration
Combined	$-23/ + 12$	$\times 0.77$ – $\times 1.12$ on every limit
Signed terms. A bias is not an error bar and is not added to one.		
Channel response	$\times 2.3$	carried by the injections, so already inside the completeness above
its residual	$+8$	the injected profile falls short of the exact response: conservative
Phase centre, annual parallax	$+1.5$	applied window by window; every limit shallower
Decorrelation	$-12/ + 17$	declared from the literature, not measured here

TABLE 6
THE RESULT IN FIVE COUNTS, EACH READ FROM THE ROWS OF TABLE 9 SO THAT THE TWO CANNOT DISAGREE. THE LINE ATTRIBUTION IS EVALUATED IN THE STELLAR FRAME AT THE ADOPTED $\pm 50 \text{ km s}^{-1}$ HALF-WIDTH, AND THE LAST TWO LINES ARE THE SPATIAL SCREEN AND THE RECURRENCE TEST APPLIED IN THAT ORDER.

Quantity	Observed
Threshold crossings	56
Line-attributed ($ \Delta v \leq 50 \text{ km s}^{-1}$, stellar frame)	19
Unattributed	37
outranking all spatial controls	1
recurring at a second epoch	0

nel $\times$ drift cell at the stellar position reaching $T \geq 5$. Such a cell is a *threshold crossing*. The molecular-line mask, $\pm 50 \text{ km s}^{-1}$ wide, evaluated in each star’s own rest frame and fixed before any crossing was classified, identifies 19 of them with a catalogued transition. 37 are unattributed by that test, and every crossing is in one column or the other. Of the unattributed, 1 outranks all 512 control positions in its own window, and it does not recur in any later epoch covering its frequency. **No crossing survives the chain, and this survey reports no detection.** Figure 8 carries the chain with the number surviving each step, Table 6 the same five counts, and Table 9 the crossing-by-crossing ledger behind them.

The mask is informative rather than permissive, and its half-width is not a tuned parameter. The width follows from the Keplerian velocity bound on circumstellar gas in these systems, and Table 10 carries every count in the chain above at half-widths from $\pm$ to $\pm \text{km s}^{-1}$: the number outranking every control is the same at all of them, and it is the same crossing. The mask attributes 19 crossings where its occupancy of the searched band would give 2.3 by chance alone, so the attributions are a kinematic result and not an accident of bandwidth.

5.3.1. Most of the crossings are expected by chance




FIG. 5.— Why the completeness is quoted in two strata. (a) The 21 directly injected Class A windows, each at the power that recovers nine carriers in ten, against the maximum of its own control ring. Marker shape is the band, open symbols the compact array, a cross a channel narrower than 200 kHz, and arrows the three windows that never reach ninety per cent. The line and band are the noise-dominated value and its spread, the dotted line the single factor fitted to all of them. (b) The two strata’s recovery curves: solid as the screen is applied, dashed with the star’s own leak into the ring removed, which is a diagnostic and not the test used here. The disc-affected pair lie on top of one another.




FIG. 6.— (a) Class A sensitivity: the cumulative number of systems whose best window reaches a given $EIRP_{90}$, the power at which nine carriers in ten are recovered through the trigger and the spatial screen together. The shaded band is the combined interval of Table 5. (b) The temporal selection function: the systems grouped by the longest separation between their searched epochs, an independent epoch being a separation of more than 1 day. **For the 27 systems observed in a single block, and the 13 whose blocks all fall within one day — left of the dashed line — independent confirmation was not possible at all, whatever the sensitivity in panel (a).** Recurrence over years can be tested for 23 systems of 82.

A 5σ trigger is a threshold on a statistic and not a false-alarm rate: the number of searched channel $\times$ drift cells per window spans several decades, so a common threshold buys very unequal protection. Summing each window’s own measured false-alarm probability over the covered Class A windows gives 34.5 expected unattributed Class A crossings, against 35 observed. *About two thirds of the crossings this survey reports are expected by chance*, which is why the chain never treats a crossing as a detection and why no artefact hypothesis is needed to account for their number.

5.3.2. One execution block accounts for four of the crossings

4 crossings belong to BD+05 1668 (GJ 273, Luyten’s Star), with the highest statistics in the survey outside

β Pictoris and HD 48370. All 4 lie in one ACA execution block, A002_Xc7fa6f_X28a8, in which the whole field is above threshold: 512 of 512 control positions exceed $T = 5$, against a survey median control maximum of 2.75. The single-integration noise is as predicted (0.1609 Jy against 0.1611 Jy), but the residual does not integrate down; a jackknife on the channel mean gives 0.030 Jy against 0.0048 Jy thermal. A slow drift common to every channel and every position therefore inflates the statistic across the field. The 24 sibling executions of the same field, with the same array and the same spectral setup, reach a maximum of $T_{\star} = 4.26$ and produce no crossing. *The median of the control ensemble, not its maximum, is the quality criterion this kind of search needs.* We report it as a diagnostic and do not adopt

it as a cut, because a cut strict enough to remove this block would also remove the HD 48370 carbon-monoxide control.

5.3.3. Stellar flares are excluded by a broadband prediction

The hosts are not a quiet sample: 2 of 2 carry activity indicators, and magnetically active stars do produce coherent bursts at centimetre wavelengths. A millimetre flare, however, raises every channel of a window by the same flux density. A flare able to produce a 5σ single-channel crossing must therefore produce a $5\sqrt{N}\sigma$ continuum event in the same integrations. Over the 48 crossings with a recoverable continuum measurement that prediction runs $49\text{--}4,717\sigma$, while the observed continuum statistic has median -0.29 and $|Z| < 3$ in 45 of 48. One block settles it directly. AU Mic A002_xa42f75_X319 contains both a crossing ($T_\star = 5.27$) and the published millimetre flare of MacGregor et al. (2020). In the crossing’s own time window the continuum is raised at $Z = 10.3$, while the crossing’s de-drifted track reaches only $Z_{\text{max}} = 2.1$. *A real, bright, independently published millimetre flare in the same data contributes 0.17σ to the single-channel statistic.*

5.4. β Pictoris: an astrophysical positive control, and the obstacle it identifies

A search that finds nothing must show that it could have found something, and the survey’s strongest features show exactly that. The frozen pipeline recovers the β Pictoris circumstellar carbon monoxide blind, without being told the disc is there. It recovers it in two rotational transitions and in 9 independent execution blocks spanning 8.4 yr: 5 Band 3 windows at $T_\star$ up to 30.76, and 4 Band 6 windows at 11.68. The identification is kinematic and does not rest on those statistics. Carried through the frame chain — to the barycentre at each block’s own epoch, then to the star’s rest frame — all 10 crossings lie within 0.5 km s^{-1} of the stellar systemic velocity, in the two lowest CO transitions of a star with mapped circumstellar CO (Dent et al. 2014; Matrà et al. 2017). The Band 3 peaks agree across blocks to 1.00 channel once each block’s own tuning is removed, so any apparent drift is far below the grid searched. The alternative reading requires a transmitter placed by coincidence at matched offsets from two CO transitions of one star.

β Pictoris is also the control that shows the recurrence test of §5.6 working in the positive direction. Every β Pictoris carbon-monoxide feature the search flagged is recovered in an independent execution block, and the screened Band 6 window returns $T_\star = 8.95$ in the second block of its own recurrence set. A non-recurrence elsewhere means nothing unless the same machinery recovers a repeat where one exists.

One of those blocks is a counterexample, and it is the more useful half of the control. In it the control ring is brighter than the star, and several of the 512 control positions exceed the stellar statistic, so the spatial screen rejects a window containing carbon monoxide that is known to be real and known to be centred on the star. The disc is resolved on the scale of the primary beam in that configuration, which is exactly the geometry the screen is built to reject. **The screen is therefore not a test of authenticity but a test of compactness,**




FIG. 7.— The β Pictoris positive control, in both directions. (a) The spatial screen. One point per execution block in which the blind search recovers circumstellar carbon monoxide, plotted as the stellar statistic against the maximum over that window’s 512 control positions: the 9 blocks in which the star outranks every control, and one further block in which the ring wins. *The control is that the search recovers a real narrow line at a stellar position, in two transitions and in every one of these blocks.* It is not that the star outranks the ring everywhere. In block X7116f1_X1785, Band 6, the control ring reaches 14.13 against 9.68 at the star, with 6 of the 512 control positions above it, so the screen rejects a window holding carbon monoxide that is real, centred on the star and independently mapped. The disc is resolved on the scale of the primary beam in that configuration, which is the geometry the screen is built to reject, so **the screen measures compactness and not authenticity** (§5.4). (b) Line confusion, which is what makes the screen’s job hard. Each crossing is plotted by its velocity offset from the nearest masked transition in its own star’s rest frame, with the adopted $\pm 50\text{ km s}^{-1}$ half-width marked. The 10 β Pictoris carbon-monoxide crossings land $0.1\text{--}0.5\text{ km s}^{-1}$ from CO, at the resolution of the systemic velocity itself, while the two CS(5 $\rightarrow$ 4) crossings of the same star sit $3765\text{--}3870\text{ km s}^{-1}$ from anything catalogued and are not attributed at any width. 39 further crossings are drawn in grey, with HD 48370 marked; 52 of the 56 are plotted, the remainder being crossings for which the release stores no control ring. *A crossing is identified by where it falls in this plane and not by how strong it is, which is why the line mask and not the trigger is the productive test.*

and on a real, spatially extended source it fails in the conservative direction. That is a stronger statement than a clean positive control would be, because it bounds the screen’s reach as well as demonstrating it (Fig. 7).

HD 48370 is the complementary case and is equally instructive. Its crossing reaches $T_\star = 26.83$, but the control ring rises with the star, to 18.36, as emission filling the primary beam does. The velocity matches the prominent CO(2 $\rightarrow$ 1) that Cataldi et al. (2023) report toward the source and flag as possible foreground cloud emission. The block’s $^{13}\text{CO}(2\rightarrow 1)$ window is bright at the ring and faint at the star. Extended line emission is therefore rejected by the same spatial statistic that promotes β Pictoris.

Taken together, these two results identify the dominant obstacle for millimetre technosignature searches, and it

is not sensitivity. The strongest apparent signals in this survey are circumstellar and foreground molecular emission, and 19 of the 56 crossings land on a catalogued transition. The mask that removes them costs per cent of the unique-frequency union over species, and it is nevertheless the single most productive test in the chain. The warning for future work is about line lists rather than about depth, and Appendix C puts a number on it.

5.5. The crossings that survive the line mask

1 of the 56 crossings is both unattributed and outranks every one of the 512 control positions in its own window: 61 Vir at 346.313811 GHz, in Band 7. The event closest to it in the ledger, CP-72 2713 at 344.269746 GHz, lies -12 km s^{-1} from SO(8₈-7₇) in the star's own frame, and the mask therefore attributes it. These are the two most interesting events in the experiment, and everything measured about them is in Table 7 rather than scattered through the text.

Both are absent when the same frequency is observed again. CP-72 2713 is the cleaner case. The next observation of that frequency began 2.05 hours later and is deeper; a carrier persisting at the discovery flux would have returned $T_\star = 6.30$ there, and the value measured at the predicted cell is 3.60. That excludes a persistent source at 7.1σ — from a repeat two hours later, with no appeal to a line list, a rank or a model. 61 Vir has 3 covering repeat blocks of the 4 the catalogue records, the nearest 1.97 hours away, expected $T_\star = 9.70$ against 2.71 measured: 9.3σ . Neither is carried forward.

5.5.1. The visibility fit, and why it is not the discriminant

51 of the 56 crossings were re-measured in the visibility domain, by phase rotating to the star and following the recovered drift (Table 8). The result sits beside the chain rather than inside it because a noise maximum at the stellar position has flux at the stellar position: at the trigger the fit returns a real part close to the trigger itself and adds little to the trigger that selected it. Injected sources in the bin the unattributed crossings occupy return $\text{Re}/\sigma = 6.58 \pm 1.97$, and the two events above are consistent with a real source and do not exclude one. **The discriminating evidence is the line mask and recurrence.**

5.6. Recurrence, applied to every unattributed crossing

A persistent transmitter should be there the next time the same frequency is observed. That is the confirmation criterion, applied in the star's own frame: the expected channel follows from the Earth's motion at the repeat epoch and the star's systemic velocity, the discovery drift is transported to it, and the survey's own statistic is read at that one cell. Nothing is maximised, so what comes back is an amplitude with an uncertainty.

The test is run on *every* unattributed crossing with any covering repeat coverage, and not only on those that outrank their controls: the control rank has no calibrated false-alarm probability, so it cannot be allowed to decide which crossings are tested. 36 of the 37 were tested, over 439 covering repeat windows in 146 execution blocks, 38 of them from the archival extension searched with the frozen pipeline after the statistic and the mask were fixed. *Nothing recurs.* The largest statistic found at any predicted cell is $T_\star = 4.71$, and 2.48 at the matched drift,

against a trigger of 5. Every tested crossing excludes a carrier persisting at its discovery flux by at least 5.0σ , median 11.6σ ; the four recurrence columns of Table 9 give the block count, the largest statistic found there at any drift, the statistic a persistent carrier would have produced and the exclusion, crossing by crossing.

One crossing could not be tested at all, and is named rather than counted: HD 14055 at 331.1559 GHz in block X1001fdc_X29ae, whose discovery dynamic spectrum was not retained, so there is no amplitude to transport into the 30 later epochs that do cover it. A further 2 blocks the catalogue lists as covering a crossing frequency have also lost their spectra and are not counted as repeat coverage. Nor is the archive exhausted: 18 further in-scope blocks carry no pipeline calibration product in the ALMA delivery at all (§3.2).

η Crv shows why the frame correction is not a detail. Its crossing at 357.281464 GHz has exactly 2 covering epochs in the archive, the second 2.19 times deeper; matching in the stellar frame moves the expected cell by -80.2 channels, so a sky-frequency test would have read the wrong channel. There the statistic is -0.618 where a persisting carrier would have given 10.95: *excluded at 11.57σ from one epoch*, carrying 4.79 equivalent discovery epochs of sensitivity.

5.6.1. What can be searched, and what can be confirmed

A sensitivity limit and a confirmation limit are different statements. Call a transmitter persistent if it would be detectable above the quoted power in every independent observation whose coverage contains its frequency. The experiment then has two extents. What it *searched* is 60 systems, 402 Class A windows and 181.5 GHz of fine-channel coverage. What it could have *confirmed* is smaller: 30 of those systems, and 79.7 GHz or 44 per cent of the Class A coverage, were observed twice at the same frequency more than a day apart; beyond a year, 11 systems and 17.8 GHz. Across the crossings tested the usable baselines run from 0.46 hours to 1.8 yr, and for 20 of 36 the nearest repeat falls within a day. **These tests exclude a carrier persistent over hours to days where there is repeat coverage; they bound no duty cycle measured in months or years, and over half the searched systems they bound nothing at all.**

5.7. A stellar-frame multi-epoch stacked search

Combining the epochs of a star in that star's own rest frame deepens the limit. 69 of the surveyed stars hold more than one searched execution block. Over 368 stacked groups, 2319 epoch-windows and 1484 h, the flux limit improves by a median factor 1.46 over each star's best single epoch (2.39 at the 90th percentile, 4.45 at most). Every stacked limit is quoted on the same completeness convention as the rest of the paper, measured by injection into the stack itself rather than read off a stacked spectrum as a threshold.

A fixed stellar-frame frequency is not a fixed observed frequency, so the epochs must be registered before they are added. The observer's motion and the star's systemic velocity move the observed frequency of a stellar-frame carrier by tens of km s^{-1} between epochs. The barycentric term alone swings by a median 8.2 km s^{-1} between

TABLE 7

THE TWO MOST INTERESTING EVENTS. EVERYTHING MEASURED ABOUT THE ONE UNATTRIBUTED CROSSING THAT OUTRANKS ALL 512 CONTROL POSITIONS, AND ABOUT THE ONE THE MASK ATTRIBUTES TO A SULPHUR MONOXIDE TRANSITION. T_* IF PERSISTENT IS WHAT A CARRIER AT THE DISCOVERY FLUX WOULD HAVE RETURNED AT THE PREDICTED STELLAR-FRAME CELL OF THE REPEAT COVERAGE COMBINED; T_* MEASURED IS THE LARGEST VALUE FOUND THERE; THE EXCLUSION IS THE DIFFERENCE IN UNITS OF THE COMBINED NOISE.

	61 Vir	CP-72 2713
ALMA band	7	7
Frequency (GHz)	346.313811	344.269746
Channel width (kHz)	488.28	488.28
Drift rate (Hz s ⁻¹)	+3182	-2056
T_*	5.958	5.809
Highest control T	5.65	5.68
Flux density (mJy)	118.8 ± 19.9	11.3 ± 2.0
EIRP (W)	5.1 × 10 ¹⁴	8.9 × 10 ¹⁴
Nearest transition	CO(3-2)	SO(8-7)
Offset (km s ⁻¹)	+457	-12
Repeat blocks	3	5
Nearest repeat (h)	1.97	2.05
Longest baseline (d)	3	2
Covering blocks in the catalogue	4	1
T_* if persistent	9.70	6.30
T_* measured	2.71	3.60
Exclusion	9.3σ	7.1σ
Disposition	not confirmed	attributed

TABLE 8

THE VISIBILITY-DOMAIN RE-MEASUREMENT AND THE THREE SETS OF CONTROL POSITIONS THE EXPERIMENT USES: THE RANK SCREEN COMPARES THE STAR WITH 512 POSITIONS IN THE IMAGE PLANE, FULL DYNAMIC SPECTRA ARE RETAINED FOR 8 OF THEM, AND 12 POSITIONS ARE REACHABLE IN THE VISIBILITIES.

Crossings re-measured in the visibility domain	51 of 56
consistent with a point source at the stellar position	41
of those, unattributed	31
Control positions in the image-plane rank screen	512
with a full dynamic spectrum retained	8
Control positions available in the visibility fit	12

the epochs of one star, and by as much as 53.9 km s⁻¹, which is many channels at the 488 kHz Class A channel width. Added as observed, the epochs of a real carrier would land in different channels and dilute rather than reinforce. Each epoch's frequency axis is therefore transformed to the stellar rest frame before stacking. The transform is applied per execution block rather than per star, because within a block the correction moves only 0.072 km s⁻¹ and a single mid-time suffices. Epochs are then divided by their own primary-beam response and combined by inverse variance over the channels of common stellar-frame coverage, with a minimum of 2 blocks and 64 common channels. Groups are keyed on the star's catalogued identity rather than on the archival directory the search ran in, and a group is the intersection of the windows covering one stellar-frame frequency, so two tunings sharing no channel cannot combine.

One kind of product is excluded before any of this. Of the 2838 retained spectra, 27 (1.0 per cent) carry archived visibility weights reaching against a survey median of , a factor , while their own time-averaged spectra are numerically zero to within a few nJy. Such an epoch would take essentially all of the weight in an inverse-variance sum; the fault is in the delivered weights, not in the stacking. Excluding them costs two-epoch groups, which fall below the 2-block minimum, and every star




FIG. 8.— What disposes of a threshold crossing. 56 crossings; 19 attributed to a catalogued transition by the stellar-frame mask; 37 unattributed, of which 36 are tested for recurrence against 439 covering repeat windows and none recurs. The control rank is an annotation and not a step: it says which crossing was examined first, not which survives. Every count is read from the ledger, so the figure and Table 9 cannot disagree.

involved retains other stacks.

The criterion was fixed before any stellar position was examined. A group is a candidate if its stacked $Z \geq 6.0$ at a common channel and above every stacked control maximum. Two further clauses guard against one epoch carrying the result: no single epoch may supply more than 0.60 of the weighted sum, and at least 0.50 of the signal must survive removal of the strongest epoch. Controls are stacked identically and every clause is evaluated at them as well. That measures the false-alarm rate of the whole test at 0.21 per cent, against the 1 per cent ceiling fixed in advance.

Registration is verified on real lines before it is used on anything else. The per-block transform collapses HD 285968's molecular line from a 19.26-channel topocentric spread to 0.798 channels, an rms of 0.351 channels = 0.223 km s⁻¹ and a factor 24 improvement, across 4 epochs spanning 44 d in 488 kHz channels. It also places β Pictoris' CO(2-1) at +3.52 km s⁻¹ in its 244 kHz windows against +3.55 km s⁻¹ in its 15.3 kHz ones. The two agree to 0.087 km s⁻¹ across a factor 16 in channel width.

The gain is $\sqrt{N_{\text{eff}}}$ and not $\sqrt{N}$, where $N_{\text{eff}} = \sigma_{\min}^2 \sum_e \sigma_e^{-2}$ is the equivalent number of equal-depth epochs. The median N_{eff}/N is 0.65 and the tenth percentile 0.30. Measured against $\sqrt{N_{\text{eff}}}$ the realised improvement is 0.992, so the combination is ideal to about one per cent. Against $\sqrt{N}$ it is 0.78, and the shortfall is unequal epoch depth rather than registration error. An epoch count alone is misleading: HD161868 stacks 13 epochs for a gain of only 1.15, because its N_{eff} is 1.7. Carrier-equivalent, the median gain is 1.27 after the measured sub-channel registration loss of $1 - 0.875$. Every stacked limit also carries a noise-realisation term of ± 7

TABLE 9

THE CROSSING LEDGER. EVERY ONE OF THE 56 THRESHOLD CROSSINGS: THE STAR, ITS EXECUTION BLOCK, THE ALMA BAND, THE CROSSING FREQUENCY AS THE CORRELATOR DELIVERED IT, THE SAME FREQUENCY IN THE STAR'S REST FRAME, THE SEARCH STATISTIC T_* , WHETHER THE STAR OUTRANKS ALL 512 CONTROL POSITIONS, THE NEAREST MASKED TRANSITION, THE STELLAR-FRAME OFFSET Δv_* FROM IT, AND THE DISPOSITION. A CROSSING IS ATTRIBUTED WHERE $|\Delta v_*| \leq 50 \text{ km s}^{-1}$; THE TRANSITION NAMED IS THE NEAREST IN THAT FRAME, SO FAR FROM EVERY TRANSITION IT IS THE MAGNITUDE AND NOT THE NAME THAT CARRIES INFORMATION. ν_* FOLLOWS FROM ν_{topo} BY THE BARYCENTRIC TERM OF THAT BLOCK'S EPOCH AND THE STAR'S SYSTEMIC VELOCITY AND NOTHING ELSE, SO THE LAST COLUMN IS REPRODUCIBLE FROM THE TWO BEFORE IT. BLOCK IDENTIFIERS OMIT THE CONSTANT A002_ PREFIX.

Star	block	B	ν_{topo} (GHz)	ν_* (GHz)	T_*	scr line	Δv_*	disposition
β Pic	Xf5d76d_Xeb8	3	115.2603	115.2711	30.765	Y	CO(1-0)	-0.2 attributed
β Pic	Xf5d76d_Xe19	3	115.2603	115.2711	27.684	Y	CO(1-0)	-0.2 attributed
HD 48370	Xc26103_X155a	6	230.5117	230.5243	26.834	Y	CO(2-1)	-17.8 attributed
β Pic	Xf4df6f_Xc7	3	115.2620	115.2711	17.907	Y	CO(1-0)	-0.2 attributed
BD+05 1668	Xc7fa6f_X28a8	6	240.1016	240.1054	17.519	n	CS(5-4)	-5911.9 unattributed; below rank screen
β Pic	Xf5d76d_Xcc5	3	115.2604	115.2711	17.288	Y	CO(1-0)	-0.2 attributed
β Pic	Xf5d76d_X32f1	3	115.2606	115.2713	14.654	Y	CO(1-0)	+0.4 attributed
BD+05 1668	Xc7fa6f_X28a8	6	224.7422	224.7458	11.868	n	¹³ CO(2-1)	+5913.1 unattributed; below rank screen
β Pic	Xd9668b_X3a90	6	230.5161	230.5378	11.681	Y	CO(2-1)	-0.3 attributed
BD+05 1668	Xc7fa6f_X28a8	6	226.7422	226.7458	10.136	n	CO(2-1)	-4931.4 unattributed; below rank screen
β Pic	X6f134f_X1484	6	230.5284	230.5379	9.832	Y	CO(2-1)	-0.1 attributed
β Pic	X7116f1_X1785	6	230.5272	230.5384	9.677	n	CO(2-1)	+0.5 attributed
β Pic	Xd9668b_Xa9df	6	230.5160	230.5377	8.948	Y	CO(2-1)	-0.4 attributed
BD+05 1668	Xc7fa6f_X28a8	6	242.1953	242.1992	8.784	n	CS(5-4)	-3349.2 unattributed; below rank screen
β Pic	Xa7a216_X23e4	6	230.5273	230.5377	7.988	Y	CO(2-1)	-0.4 attributed
61 Vir	Xc079b5_X82f	7	346.3138	346.3232	5.958	Y	CO(3-2)	+457.1 unattributed; rank-flagged
HD 285968	Xd98580_X3bf6	6	230.5021	230.5447	5.956	n	CO(2-1)	+8.7 attributed
HD 285968	Xdb7ab7_X5b4e	6	230.5116	230.5447	5.868	n	CO(2-1)	+8.7 attributed
HD 285968	Xdb4b9a_X108f	6	230.5100	230.5447	5.835	n	CO(2-1)	+8.7 attributed
CP-72 2713	Xff0235_X4a6d	7	344.2697	344.2965	5.809	Y	SO(8 ₈ -7 ₇)	-12.3 attributed
HD 139084 B	Xcd8029_Xb6b0	7	345.8241	345.8241	5.763	n	CO(3-2)	+24.4 attributed
HD 14055	Xff3d8_Xd46	7	331.7699	331.7664	5.575	n	SO(1 ₂ -0 ₁)	+2167.0 unattributed; below rank screen
HD 23484	X10b22d7_X2445	6	230.7904	230.8006	5.520	n	CO(2-1)	+341.5 unattributed; below rank screen
HD 161868	Xff45a6_X10ca8	7	345.6338	345.6491	5.481	n	CO(3-2)	-127.4 unattributed; below rank screen
HIP 10679	Xa95c04_X14cd	6	220.4185	220.4035	5.456	n	¹³ CO(2-1)	+6.5 attributed
HD 14055	X11adad7_X25c5e	7	345.4564	345.4303	5.448	n	CO(3-2)	-317.0 unattributed; below rank screen
HD 207129	Xc19d6f_X5706	6	230.4226	230.4070	5.433	n	CO(2-1)	-170.3 unattributed; below rank screen
β Pic	Xd9668b_X3a90	6	241.7512	241.7739	5.420	n	CS(5-4)	-3869.7 unattributed; below rank screen
HR 1010	Xc6fa08_X1053	6	230.5133	230.5287	5.410	n	CO(2-1)	-12.0 attributed
ALMA J153702653-33192492	Xff0235_X8392	6	230.5222	230.5392	5.405	n	CO(2-1)	+1.6 attributed
HD 110411	Xe45e29_Xb7d4	6	219.4538	219.4357	5.393	n	C ¹⁸ O(2-1)	-170.2 unattributed; below rank screen
HD 131392	X129754b_X6946	6	231.1998	231.2231	5.289	n	H30 α	-876.3 unattributed; below rank screen
HD 14055	Xff99e1_X24be0	7	331.7078	331.7035	5.277	n	SO(1 ₂ -0 ₁)	+2109.7 unattributed; below rank screen
AU Mic	Xa42f75_X319	6	230.6864	230.6705	5.274	n	CO(2-1)	+172.3 unattributed; below rank screen
HD 69830	X12209f7_X16edb	7	322.4325	322.4444	5.259	n	SO(1 ₂ -0 ₁)	-6317.5 unattributed; below rank screen
ALMA J153702653-33192492	Xfe62c1_X398	6	216.1844	216.2036	5.228	n	SiO(5-4)	-1244.7 unattributed; below rank screen
AU Mic	X899169_X1838	6	230.8252	230.8294	5.221	n	CO(2-1)	+379.0 unattributed; below rank screen
HD 14055	Xff99e1_X2c3bb	7	345.5256	345.5216	5.186	n	CO(3-2)	-237.9 unattributed; below rank screen
HD 14055	Xff99e1_X3835	7	345.0099	345.0070	5.177	n	SO(8 ₈ -7 ₇)	+606.4 unattributed; below rank screen
HD 14055	Xff99e1_X2b2ee	7	344.7007	344.6965	5.175	n	SO(8 ₈ -7 ₇)	+336.0 unattributed; below rank screen
LHS 1140	Xd9d7c7_X6fd2	6	230.6355	230.6268	5.136	n	CO(2-1)	+115.5 unattributed; below rank screen
HD 10647	Xff294f_Xaf70	7	331.6570	331.6943	5.120	n	SO(1 ₂ -0 ₁)	+2101.4 unattributed; below rank screen
CD-57 1054	Xcf525f_X5aa4	7	346.5099	346.5265	5.108	n	CO(3-2)	+633.4 unattributed; below rank screen
ALMA J153702653-33192492	Xf8f6a9_X10359	6	217.7292	217.7269	5.091	n	H ₂ CO(3-2)	-680.5 unattributed; below rank screen
GJ 849	Xd9668b_X841a	6	230.8769	230.8581	5.090	n	CO(2-1)	+416.3 unattributed; below rank screen
ALMA J153702653-33192492	Xfe62c1_X398	6	218.8729	218.8923	5.090	n	C ¹⁸ O(2-1)	-912.2 unattributed; below rank screen
HD 14055	X1001fdC_X29ae	7	331.1559	331.1546	5.082	n	SO(1 ₂ -0 ₁)	+1610.2 unattributed; below rank screen
β Pic	Xd9668b_Xa9df	6	241.8367	241.8594	5.079	n	CS(5-4)	-3765.1 unattributed; below rank screen
ALMA J153702653-33192492	Xf8f6a9_X10359	6	218.6711	218.6688	5.048	n	H ₂ CO(3-2)	+613.5 unattributed; below rank screen
HD 69830	X12277ed_X178c6	7	321.2005	321.2149	5.046	n	SO(1 ₂ -0 ₁)	-7436.5 unattributed; below rank screen
HR 1010	Xc66ce1_Xaa9	6	230.8348	230.8491	5.044	n	CO(2-1)	+404.6 unattributed; below rank screen
61 Vir	Xc067f7_X822	7	345.5599	345.5678	5.038	n	CO(3-2)	-197.8 unattributed; below rank screen
HD 170773	Xff0235_X8c42	7	331.0101	331.0234	5.029	n	SO(1 ₂ -0 ₁)	+1490.8 unattributed; below rank screen
ALMA J153702653-33192492	Xd9d7c7_Xbb64	7	347.2179	347.1890	5.009	n	CO(3-2)	+1207.7 unattributed; below rank screen
HD 14055	X1190de8_Xd526	7	344.4676	344.4515	5.006	n	SO(8 ₈ -7 ₇)	+122.7 unattributed; below rank screen
η Crv	X122b6ff_X1041e	7	357.2815	357.2475	5.006	n	CO(3-2)	+9928.0 unattributed; below rank screen

per cent, because the noise scale is a median absolute deviation over a median of 127 common channels.

The deepest stacked limit is Proxima Centauri's, and it is a Class B limit. Proxima holds 66 blocks over 85 h spanning 6.3 yr, all of them coarse 15.625 MHz windows, so the stack constrains an unresolved spectral excess and discriminates no drift. Its N_{eff} is only 7.9, and it is the second number that sets the realised gain of $3.19\times$ over the same star's best single epoch. Carriers injected into every epoch of that stack and recovered through the same path give a 90 per cent recovery point at $3.84\times$ the stack's nominal threshold. On that convention the best single epoch reaches W and the stack $\text{EIRP}_{90} = 1.73 \times 10^{13} \text{ W}$, for an emitter steady in the stellar frame to 20.9 km s^{-1} — one 15.625 MHz channel — over epochs spanning 6.3 yr, with no inter-epoch drift above Hz s^{-1} . The gain, not the absolute figure, is what stacking bought: at pc the depth is mostly geometry. The calibration factor is a property of the stacking chain rather than of this star, and a second group returns $3.56\times$.

2 groups pass the full criterion and both are known lines. They are β Pictoris' CO $J = 1-0$ at $+3.09 \text{ km s}^{-1}$

over 8 epochs ($Z = 45.3$) and its $J = 2-1$ at $+3.55 \text{ km s}^{-1}$ ($Z = 12.1$). The two bands agree to 0.46 km s^{-1} , or 0.73 of a CO(1-0) channel. That is the astrophysical positive control passing in the stack. Attributing every survivor against the survey's own $\pm 50 \text{ km s}^{-1}$ mask leaves 0 unattributed. The stack is also an independent test of the threshold crossings of §5.6: the 9 stars carrying them stack to $Z = 2.82-4.43$ in 8 cases, 0 exceed their own control ensembles and 0 passes the criterion. The exception is HD 285968, which reaches $Z = 12.89$ on the foreground line used above as the registration standard and is rejected because its own controls reach 14.19: the emission is extended, as a cloud is and a transmitter is not.

The restriction travels with every number above, and its width is the channel. At CO(2-1) one 488 kHz channel is km s^{-1} , so the $\pm 30 \text{ km s}^{-1}$ reflex swing of a transmitter on a short-period giant planet, 23.1 MHz, spans 47 of them and de-registers itself completely between epochs — but only channels of a 15.625 MHz window, which is why the coarse stacks tolerate reflex motion and buy that tolerance with bandwidth. Intrinsic drift divides

the same way: a stack in 488 kHz channels is confined to $1.3 \times 10^{-1} \text{ Hz s}^{-1}$ over the median inter-epoch span of 44 d and $2.0 \times 10^{-3} \text{ Hz s}^{-1}$ over the longest, 2868 d. Each stacked limit therefore constrains persistent emitters that are stable in the stellar frame to within its own channel, and nothing else. The single-epoch limits of §6.3 are the ones that apply generally.

6. DISCUSSION

6.1. *What is different about a millimetre technosignature search*

The most transferable result of this experiment is not the non-detection. It is a measurement of what actually limits a carrier search at millimetre and submillimetre wavelengths. The binding constraints are molecular-line confusion, a channelisation chosen by other observers for other purposes, and when the archive happened to look. Thermal sensitivity, which is what usually limits a centimetre-wave search, is not among them: ALMA supplies more of it than this experiment can use.

Confusion, and not interference, sets the false-positive rate. Catalogued molecular transitions, almost all of them carbon monoxide, account for 19 of the 56 threshold crossings, and no crossing in this survey is attributed to interference (Appendix E). 10 of the 12 windows that reached the spatial stage are disc or foreground carbon monoxide, arising in 2 of the 69 debris-disc systems searched. The remedy that works is a velocity-space mask evaluated in the stellar frame, and it is not free: it removes 4.1 per cent of otherwise usable coverage.

Spatially extended emission is the second discriminant, and it has no real analogue at centimetre wavelengths. A line that fills the primary beam raises the star and the control annulus together instead of cancelling, so the spatial screen loses power exactly where circumstellar gas is brightest. Two targets show both faces of this. β Pictoris is the benign one: the frozen pipeline recovers its circumstellar carbon monoxide blind, in two transitions, which is the positive control the experiment needed. HD 48370 is the awkward one. Its resolved disc emission lies close to a band edge, and that single window contributes 59 per cent of all the near-edge threshold exceedances in the survey. Pooled over the searched windows the near-edge exceedance rate is $\times 2.464$ the interior rate; drop that one window and it is $\times 1.020$, and on the half of the windows held out from its own stratum it is $\times 1.181$ ($p = 0.24$). What looked like a property of the band edges is one resolved disc, so no edge correction is applied to any limit.

The channelisation is inherited, and it decides what kind of search is possible at all. 1249 of the 1651 windows searched have channels so coarse that every drift in the grid stays within one channel over a whole track; they measure an amplitude but cannot tell a drifting carrier from a stationary one, which is why the carrier search rests on 402 of them. Across the archive the native channel widths span about seven orders of magnitude (Fig. 1(b)), and nothing in the search could choose them.

Repeat epochs are what turn a crossing into a candidate, and an archive supplies them only incidentally. Temporal sampling is the weakest axis of this experiment, and for close to half the sample it constrains an intermittent emitter hardly at all, whatever its power.

Coverage is also incomplete at the level of whole observations: 144 of the 177 execution blocks in scope were searched, 18 carry no observatory calibration at all, 11 failed calibration, 1 is infeasible at this data volume and 3 were not attempted. For one of them the observatory delivered raw data it had never calibrated, so part of that gap lies in the delivery rather than in the processing.

Finally, a threshold is not a sensitivity, and here the difference is large. A nominal 5σ level is a property of the noise, not a statement about what a transmitter of a given power does to a pipeline. Measured by injection and recovery through the complete automated selection, the power needed for 90 per cent recovery is 4.42 times the trigger power in Class A and 4.70 in Class B. One limitation should be read with every number here: every check reported in this paper tests this pipeline against itself, and because the only other archival search at these frequencies targets a disjoint sample, there is no external pipeline to compare against.

6.2. *The search space this archive opens*

The carrier search is 402 fine-channel windows toward 60 stellar systems, covering 47.7 GHz of unique sky frequency; the 1249 coarse-channel windows are a supplementary spectral-excess experiment that cannot search for drift at all. Three statements about novelty are worth keeping apart, because they are easily run together. No previous search has looked for technosignatures toward stars within 40 pc at any frequency above about 30 GHz. One previous search has worked in ALMA Band 3, but toward stars at 1.01 kpc and beyond (Mason et al. 2025), so Band 3 is not searched toward nearby stars either. And 20.6 GHz of the frequency covered here lies inside the 10 MHz–115 GHz interval that Wright et al. (2018) adopt as the frequency axis of the search haystack, which is a definition of parameter space rather than coverage by any telescope.

Bandwidth alone overstates what that buys, because different frequencies were searched toward different numbers of stars and to very different limits. Figure 9 therefore replaces 47.7 GHz as the description of coverage with the quantity a reader needs: how many independent systems would have yielded a recovery at 90 per cent completeness, as a function of sky frequency and transmitter power. At an EIRP of W the covered frequency falls to GHz, per cent of the union, reaching systems in all and at the best-covered frequency near GHz. With no limit on power that same frequency reaches of the 60 systems, and half the sample is reached only above W .

Archival SETI buys sky and frequency cheaply and buys confirmation not at all. Across both classes this experiment reached 82 systems and 118.1 GHz without an allocation, at a sensitivity no targeted programme at these frequencies has matched, on data selected by others. 90 per cent of the searched stars were observed for disc or planet-formation science, which puts circumstellar carbon monoxide in a large fraction of the fields, and only 21 of the 5908 catalogued M dwarfs within 40 pc have any coverage at all. The quantities here describe *these* stars, frequencies and epochs and nothing wider.

For comparison with other programmes the conventional measure is the continuous-waveform transmitter figure of merit of Enriquez et al. (2017), $\text{CWTFM} = \zeta_{\text{AO}} \text{EIRP}_{\text{min}} / (N_{\star} \nu_{\text{rel}})$, in which $\nu_{\text{rel}} = \Delta\nu / \nu_{\text{mid}}$ is the




FIG. 9.— Where the experiment could have found something. Shading gives the number of independent systems that would have yielded a recovery at 90 per cent completeness for an unresolved carrier of the plotted EIRP at the plotted sky frequency; unshaded means none. One panel per populated ALMA band, with panel widths in proportion to the frequency each contributes: the 47.7 GHz of Class A coverage is disjoint islands spread across 114.0–873.1 GHz, so a single continuous abscissa would draw every island as a hairline. The dashed line is W , at which GHz of the union survives. Figure 3 shows where ALMA looked; this shows where the experiment was sensitive.

fractional bandwidth and ζ_{AO} is fixed so that the merit is unity for 1000 stars at $\nu_{\text{rel}} = 0.5$ and an EIRP of 1×10^{13} W; smaller is better. For the Class A experiment $N_{\star} = 60$ and $\nu_{\text{rel}} = 0.097$, and taking EIRP_{min} as the largest in the sample, which is the convention, the value is 4.6×10^5 — on the median rather than the worst system, 1.9×10^4 . Enriquez et al. (2017) report 0.85, and Margot et al. (2023) correspond to 0.11, evaluated here from their published sample size, band and power limit because that paper reports a modified Drake figure of merit instead. Nearly all of the difference is the product $N_{\star}\nu_{\text{rel}}$: 60 systems across a fractional bandwidth of 0.097 is a far smaller haystack than a wideband centimetre survey of thousands of stars, and the Drake figure of merit is dominated by instantaneous bandwidth in the same way. Both compare surveys; neither is an occurrence rate, and the non-detection here cannot be converted into one because an archive-selected sample does not supply the denominator.

6.3. What this experiment excludes, and what it does not

Excluded. Consider a continuously transmitting, unresolved carrier that lies outside the molecular-line mask, inside the searched frequency range and inside the linear-drift grid, and that is present during one of the archival epochs. Such a carrier would have been found nine times in ten above the EIRP_{90} of §5.1, a median of 1.77×10^{15} W per system on its best window. None is present toward any of the 60 systems the carrier search covers.

Four further axes bound the excluded population as sharply as power does. *Frequency:* the coverage is sparse and inherited, and Fig. 9 is the statement of it; the median system is covered over only 1.84 GHz. *Drift and acceleration:* the searched grid reaches a line-of-sight acceleration of $3.60\text{--}4.00 \text{ ms}^{-2}$, which contains every known planet of the sample but one (Fig. 4). A transmitter whose own oscillator drifts faster than that kinematic budget lies outside the grid. *Persistence:* the carrier must be radiating during one of the archival observations, so these are limits on persistent emission. *Spectral*

width: the statistic is built for emission confined to one 15.3–1953 kHz channel. A carrier broadened beyond a channel loses signal to its own width, and a signal broad enough to resemble continuum is removed with the spectral baseline.

Not excluded. A transmitter absent during the archival observation, or on for less than about a tenth of the time. At a duty cycle of one tenth, recovery falls to 71 per cent for the fine class and 83 for the coarse; those trials are undrifted, so the joint drifting-and-intermittent case is not measured at all. Emission inside the $\pm 50 \text{ km s}^{-1}$ molecular-line mask, at any strength, which costs 4.8 GHz of coverage. Drift outside the searched grid, or drift curved enough to leave a channel within the track. Signals broad enough to be removed with the continuum baseline. Anything below a window’s own completeness. And any emission from the 17 566 stars within 40 pc that ALMA has not observed, which is almost all of them.

An illustrative engineering benchmark. A limit of 10^{15} W is hard to picture, so it is worth naming one transmitter that would produce it, provided the example is not mistaken for a privileged scale. A 12 m dish radiating 1 MW at 230 GHz at an aperture efficiency of 0.7 has a directive gain $G = \eta(\pi D\nu/c)^2 = 5.9 \times 10^8$ and so an EIRP of 5.9×10^{14} W, a depth this survey reaches for 14 of the 60 Class A systems. Nothing but convenience selects those numbers — aperture, efficiency, pointing, duty cycle and the chance that the beam contains the Solar System all enter, and none is constrained here — so the results of this paper remain the model-independent EIRP_{90} limits. What the example does make plain is that EIRP is not transmitter power: the same megawatt radiated isotropically would be 10^6 W of EIRP, and it is the directivity that lifts a modest transmitter into the range this survey probes. These limits therefore apply to beamed transmitters whose beam contains the Solar System, and say almost nothing about omnidirectional beacons.

Two blind spots. The search is Stokes I only, so the discrimination that polarisation would offer against unpolarised astrophysical emission is unavailable; ALMA’s feeds are linear, and the circularity of a circularly polarised carrier lies in the cross-hands, which this analysis does not extract. And the frequency coverage is the archive’s and not a choice: none of the “magic frequency” arguments that have guided centimetre-band searches falls in an ALMA band, so a transmitter deliberately placed on one lies outside this experiment by construction.

6.4. What a dedicated ALMA technosignature experiment should do differently

A purpose-built experiment on the same telescope would be far stronger, and the changes it needs follow from the limitations above. *Channelise finely everywhere:* the 1249 coarse-channel windows measure an amplitude but cannot distinguish a drifting carrier from a stationary one. *Choose the targets and the frequencies:* an archive weighted toward dusty discs is the worst case for molecular-line confusion, the M dwarfs that dominate the nearby habitable-zone planet census are the least covered, and windows placed deliberately away from the bright transitions would recover most of the 4.8 GHz the mask costs — all three are decisions a proposal can sim-

ply reverse. *Repeat the epochs on purpose:* two or three visits separated by days and by years turn a crossing into a testable claim, and the separation is a scheduling constraint rather than extra sensitivity. *Keep the controls and the polarisation:* full dynamic spectra survive for only 8 of the 512 control positions here, which is what limits the second-stage null in a structured field, and only a cross-hand extraction can see a circularly polarised carrier; both cost storage and nothing else. *Retain the visibilities:* rotating the phase centre to the star leaves a genuine point source entirely in the real part of the visibility, with an imaginary part consistent with zero on every baseline. That asks whether the data contain a source *where the star is* — a physically specific question, and a stronger one than any rank against neighbouring positions in an image.

7. CONCLUSIONS

1. ALMA’s public archive opens a region of the technosignature search space that had not been examined toward nearby stars. Of the 118.1 GHz of unique sky frequency searched here, only 20.6 GHz lies inside the frequency interval covered by centimetre-wave programmes.
2. The primary experiment is 402 fine-channel spectral windows toward 60 stellar systems, covering 47.7 GHz of unique frequency, searched for an unresolved carrier whose frequency drifts. A secondary coarse-channel spectral-excess search covers a further 70.4 GHz in 1249 windows but cannot discriminate drift, and both come from one uniform reduction of 484 public execution blocks, 4.1 TB of raw data.
3. The sensitivity is measured by injection and recovery through the complete automated selection rather than inferred from a threshold: $\text{EIRP}_{90} = 4.42 P_{\text{trig}}$, a median of 1.77×10^{15} W per system. That is deep enough for 14 of the 60 Class A systems to exclude a 12 m dish radiating 1 MW at 230 GHz, whose EIRP is 5.9×10^{14} W.
4. 56 threshold crossings are reported with their dispositions, and 19 are attributable to catalogued molecular transitions, almost all of them circumstellar or foreground carbon monoxide, among them the β Pictoris debris disc recovered blind in two transitions. None of the remaining 37 is independently confirmed: none recurs in the repeat observations available, and 1 also outrank all 512 spatial controls.
5. Combining the epochs of each star in that star’s rest frame deepens the flux limit by a median factor 1.46 over the best single epoch. For Proxima Centauri the stacked limit, measured by injection into the stack, is 1.73×10^{13} W for an emitter steady in the stellar frame to 20.9 km s^{-1} across epochs spanning 6.3 yr.
6. These results constrain persistent, unresolved carriers within the frequency, drift, duty-cycle and sensitivity domain actually searched, toward the stars ALMA happened to observe. They cannot be converted into an occurrence rate for transmitters in the 40 pc stellar population, which would require a denominator an archive-selected sample does not supply.

AUTHOR CONTRIBUTIONS

G.J.W. conceived the survey, defined the target sample and the selection criteria, and wrote the manuscript. R.D. built and operated the archival retrieval, calibration and search pipeline, implemented the spatial control statistic and the injection–recovery campaigns, and produced the released catalogue and its regeneration scripts. Both authors designed the candidate vetting procedure, carried out the statistical validation reported in §4.1 and Appendix D, and agreed the dispositions of Table 9.

FUNDING

This work received no dedicated funding. It used public archival data and no observing time was awarded for it.

COMPETING INTERESTS

The authors declare no competing interests. VBRL Holdings is a privately held company through which R.D. conducts independent research, named here solely as his affiliation; it had no role in the design of the survey, the selection of targets, the analysis, the decision to publish or the content of this paper, and provided no funding for the work.

ACKNOWLEDGEMENTS

This paper uses ALMA data from the public ALMA Science Archive. The 404 execution blocks searched span 2013 October 6 to 2025 June 11 across 63 proposal codes, which in the standard ADS/JAO.ALMA#*code* form are: 2012.1.00142.S, 2012.1.00198.S, 2012.1.00268.S, 2012.1.00385.S, 2012.1.00437.S, 2012.1.00993.S, 2013.1.00359.S, 2013.1.00588.S, 2013.1.01147.S, 2013.1.01214.S, 2015.1.00307.S, 2015.1.00633.S, 2015.1.00783.S, 2015.1.01015.S, 2015.1.01260.S, 2016.1.00104.S, 2016.1.00358.S, 2016.1.00515.S, 2016.1.00637.S, 2016.1.00803.S, 2016.1.00817.S, 2016.1.00880.S, 2016.2.00015.S, 2016.2.00200.S, 2016.A.00013.S, 2017.1.00167.S, 2017.1.00200.S, 2017.1.00215.S, 2017.1.00561.S, 2017.1.00698.S, 2017.1.00704.S, 2017.1.00786.S, 2017.1.00986.S, 2017.1.01583.S, 2017.1.01644.S, 2018.1.00045.S, 2018.1.00072.S, 2018.1.00126.S, 2018.1.00461.S, 2018.1.00470.S, 2018.1.00947.S, 2018.1.01149.S, 2018.1.01545.S, 2018.1.01833.S, 2019.1.00189.S, 2019.1.01443.T, 2019.1.01517.S, 2019.1.01603.S, 2019.1.01681.S, 2019.2.00208.S, 2021.1.00485.S, 2021.1.01209.S, 2021.2.00097.S, 2022.1.00134.S, 2022.1.00338.L, 2022.1.00793.S, 2022.1.01163.S, 2023.1.00390.S, 2024.1.00165.S, 2024.1.00595.S, 2024.1.00722.S, 2024.1.01095.S, 2024.A.00040.S. Per-block epochs are in the machine-readable release. The archive metadata exposes no proposal code for 2 of the 404 blocks (A002_X11d9ce7_X5257, A002_Xbf792a_X26ec), which are therefore not represented in that list. ALMA is a partnership of ESO (representing its member states), NSF (USA) and NINS (Japan), together with NRC (Canada), NSTC and ASIAA (Taiwan), and KASI (Republic of Korea), in cooperation with the Republic of Chile. The Joint ALMA Observatory is operated by ESO, AUI/NRAO and NAOJ. This work has made use of data from the European Space Agency (ESA) mission *Gaia*, processed by the *Gaia* Data Processing and Analysis Consortium (DPAC). This research has made use of the SIMBAD database (CDS, Strasbourg, France).

DATA AVAILABILITY

The release carries every reduced data product, the full target-selection and analysis pipeline, the 1651-row per-window catalogue, the injection records and the control statistics. Deposited beside the catalogue are the per-window search statistics, one row each, and the frozen inputs of every derived quantity in this paper: the archival harvest, the matched recurrence fits, the frequency-clustering computation, the block-by-block archival verdicts, and the ALMA delivery record for every block in scope that could not be searched. What has been demonstrated rather than claimed is narrower, and it is a measured number in both cases: 120 generated products regenerate byte-identically from the frozen inputs, and 41 headline quantities are re-derived from the released catalogue alone, without the pipeline. *The reference of record is the Zenodo deposit*, which is frozen, carries a DOI minted on acceptance, and is the version any reproduction should cite; the repository at <https://github.com/Tilanthi/SETI> is a working mirror and may move on.

SOFTWARE

The pipeline depends on CASA 6.7.6 (casatools/casatasks 6.7.6-14; CASA Team et al. 2022), Python 3 with astropy (Astropy Collaboration et al. 2022), numpy (Harris et al. 2020) and scipy (Virtanen et al. 2020), and the ALMA Science Archive and SIMBAD TAP interfaces through astroquery (Ginsburg et al. 2019); all are free and open-source, with exact versions recorded in the frozen release.

APPENDIX

SPECTRAL RESPONSE, FREQUENCY CONVENTIONS AND THE PRIMARY-BEAM RESPONSE

A reader checking a threshold in this paper needs three things: the frame every frequency is quoted in, the instrumental response that separates a channel threshold from a transmitter power, and the primary-beam factor that converts a measured noise into a flux at the star. They are collected here. The per-window catalogue, `catalogue/windows.csv` in the deposit, carries one row per searched window and the columns each of these corrections uses.

Frequency frames, linewidth and smearing

Every frequency axis here — searched spectra, tabulated ranges and crossing frequencies alike — is *topocentric* as delivered by the ALMA correlator, with no velocity-frame shift applied at any stage. Velocities enter only in the vetting stage’s line mask, which is evaluated in the same observed sky frame (Appendix C): all 1396 released `line_offset_MHz` values reproduce topocentrically against the frozen rest-frequency list. The tolerance is wide enough to absorb the barycentric and systemic terms it does not apply, so the residual error is the catalogue radial-velocity uncertainty alone, a few km s^{-1} within 40 pc; second-order relativistic terms at the 50 km s^{-1} scale are $\lesssim 5 \text{ kHz}$ at 350 GHz, below the finest native channel searched. The stacking is drift-aligned and inverse-variance weighted: coherent in frequency trajectory, not in electric-field phase.

The native channel width is also the linewidth convention for every EIRP threshold. Channel dilution spreads a sub-channel carrier’s total flux across the channel, so $\text{EIRP}_{\text{trig}}$ is the minimum *total* power of a channel-confined carrier, independent of the transmitter’s own width below the channel, and never a spectral power density. The Hz-resolution surveys’ limits are total powers in the same sense, each the best case of its own channelisation, which makes the axes of Fig. 1(a) commensurable. One number fixes the scale, illustratively and not achievably: the median Class A channel is 488 kHz, so rechannelising the same data to the 2.79 Hz of a dedicated narrowband search would, on $P_{\text{min}} \propto \sqrt{\Delta\nu}$, divide the threshold by 418. ALMA cannot do this: the correlator delivered the channelisation the archive holds, and no reprocessing recovers resolution that was never recorded.

Channel widths come from the measurement sets’ own spectral window tables (CHAN_WIDTH). A missing “EFFECTIVE_BANDWIDTH” column is not evidence of an absence of online spectral averaging, MS v2 defining no such column, so the quoted channelisations are not claimed to be noise bandwidths.

ALMA applies online Hanning smoothing by default, so the spectral response is ~ 2 channels wide and neighbouring channels are correlated (Fig. 10). For the 0.25/0.5/0.25 kernel the coefficients are $\rho_1 = 2/3 = 0.667$ and $\rho_2 = 0.167$, with $\sigma_{\text{smoothed}}/\sigma_{\text{raw}} = 0.612$: the correlation is not confined to adjacent channels, the lag-two term being a sixth. The peak-channel fractions 0.42–0.50, median 0.48, and the resulting $\times 2.00$ – $\times 2.36$ optimism of a nominal threshold are given in §4; the released catalogue carries the nominal threshold and both corrected values for every window.

The injected carriers are deposited with the correlator’s own channel response, at a sub-channel phase drawn afresh for each tone, so the smoothing loss sits inside the measured completeness rather than being applied to it. The deposited profile is the tone divided between the two channels it straddles and then passed through the three-point kernel, which leaves 0.500 of the flux in the peak channel for a carrier at a channel centre and 0.375 for one on a boundary, against 0.4244 for the exact transform of the lag window; averaged over phase the deposition therefore falls 8 per cent short of the ideal response, and the completeness is conservative by that much. Where in a channel a carrier was placed makes no difference to its recovered signal-to-noise: the medians over 3673 tones agree across the whole range of sub-channel phase to 1 per cent, because a carrier at a typical searched drift rate crosses about 11 channels during one observation and so samples the whole response. Since the loss is carried end to end it must not be applied to the result as well: at the ninety per cent point of the recovery curve the peak channel holds $1.93\times$ the nominal trigger flux, which is that factor of about two, already spent.

Two measurements fix the smoothing state. The retained per-integration residuals of the second-stage local null have a lag-one spectral autocorrelation of 0.65 against the Hanning prediction 0.667, so the correlation is instrumental and those spectra are smoothed and unaveraged; and the archive’s reported effective spectral resolution is exactly twice the channel separation in 1505 of the 1651 windows, in both correlator modes. The other




FIG. 10.— Why the nominal channel threshold is not the transmitter threshold. ALMA applies online Hanning smoothing by default, so an intrinsically unresolved carrier is spread over neighbouring channels and the peak channel holds only part of its power: half at a channel centre, 0.42 on a boundary (first three panels; blue is the peak channel, the dashed line the carrier). The search threshold is set in units of the channel noise, so the power a transmitter must radiate to reach it is larger than the nominal figure by the reciprocal of that fraction (right-hand panel). This is the whole of the difference between P_{trig} and P_{eff} .

146 report a smaller ratio, the signature of online channel averaging, where the peak-channel loss is smaller, and the 0 Band 9 and 10 windows absent from the metadata snapshot take ALMA’s default, so the blanket factor is right for the great majority of both classes. *The correlation model is truncated at lag two*: the baseline step, a per-integration median in blocks of up to 65 channels, could imprint correlation on longer scales that the released products retain no residual spectra to measure, which would make the independent-trials count of Appendix D an over-estimate.

Linewidth convention. The statistic is the single-channel flux density after drift correction, so the quoted $\text{EIRP}_{\text{trig}}$ applies to any transmitter narrower than the native channel and is flat in the assumed linewidth. Flatness is exact only for a centred tone; a boundary-straddling one divides its power between two channels, which is the loss the injections carry and the reason EIRP values are quoted to two significant figures.

Transmitters wider than one channel. The baseline filter (§4.1) is flat for features narrower than roughly half its effective width W and falls to zero above W , so the survey’s sensitivity is to channel-confined morphologies and W is the linewidth ceiling: 65–74 channels over the fine class and 30–59 over the coarse one, that is ≈ 32 MHz at 488 kHz and ≈ 0.9 GHz at 15.62 MHz. Within that ceiling, a top-hat transmitter spanning N_{ch} native channels has per-channel flux the total divided by N_{ch} at unchanged noise, so to first order

$$\text{EIRP}_{\text{trig}}(N_{\text{ch}}) \simeq N_{\text{ch}} \times \text{EIRP}_{\text{trig}}(1), \quad N_{\text{ch}} \lesssim W/2, \quad (\text{A1})$$

the rescaling the main text applies to signals broader than one channel.

Intra-integration smearing. The frequency excursion at the maximum searched drift rate across one integration is at most 1.10 channels over the 1651 windows, and a drift smeared by Δ channels within an integration keeps $\eta_{\text{smear}} = \text{sinc}(\pi\Delta/2)$ of the peak amplitude, so the effective threshold is $\text{EIRP}_{\text{trig}}/\eta_{\text{smear}}$. Only 9 windows have $\eta_{\text{smear}} < 0.99$ — HIP 10679 Band 6 (0.574, $\times 1.74$), β Pic Band 6 (0.574, $\times 1.74$), HD 172555 Band 6 (0.574, $\times 1.74$), HIP 10679 Band 6 (0.606, $\times 1.65$), HD 172555 Band 6 (0.606, $\times 1.65$), η Crv Band 8 (0.865, $\times 1.16$), η Crv Band 7 (0.932, $\times 1.07$), HD 48370 Band 6 (0.979,

$\times 1.02$) and CP–72 2713 Band 6 (0.979, $\times 1.02$) — and the remaining 442 agree with their nominal thresholds to better than one per cent. Compounding the worst smearing penalty with the worst Hanning factor gives $\times 4.6$ for those few, a product no released column carries.

The primary-beam response, and the two columns that let a reader recompute it

The sensitivity quoted for a window is $S_{\text{min}} = 5\sigma/A$, where A is the primary-beam response at the star’s offset from the phase centre: the noise σ is measured in apparent flux, the limit is quoted in true flux, and the two differ by A . That offset is not a nuisance. Archival fields are pointed at a catalogue position, often decades old, while these are high-proper-motion stars, so the star can sit well off axis: over the 1651 released windows the offset has a median of $0.089''$, a 99th percentile of $7.43''$ and a maximum of $8.13''$, and A runs from 0.999984 in the median down to 0.5522, below 0.99 on 91 windows and below 0.9 on 34. *None* falls below one half, and every one of the low-response windows is a physically explicable J2000 pointing on a fast-moving star rather than an error. A window is kept only where the response is at least 0.5, a correction of at most $\times 2$, since at low response a small error in the beam model becomes a large error in inferred flux. The offsets are bimodal and the floor sits in the gap: the last retained window has response 0.55 at $0.46\theta_{\text{PB}}$, the first window beyond it 0.35 at $0.62\theta_{\text{PB}}$, and nothing lies between, so the rule binds prospectively and excludes none of the 1651 released windows. Beyond it are the 12 ϵ Eri Band 6 windows at 1.0 – $3.5\times$, also past the first sidelobe transition where the Gaussian form fails; the two criteria are independent and select the same windows. They are *withheld*, enter no number here and are absent from the released catalogue, although under the nominal rule they would have covered 3.7×10^{13} – 1.1×10^{15} W at 3.22 pc.

The correction is applied exactly once, and that is the complete correction. The offset is computed for every measurement set from its own `FIELD::PHASE_DIR` and the propagated stellar position, so no window carries a zero-by-assumption offset. It enters the published numbers in one place only: the identity $S_{\text{min}} = 5\sigma/A$ closes on all 1651 released rows to $8.7\text{e-}05$ in relative terms. *The count is given in two parts, because on part of it the*

identity cannot fail: on 10 rows whose products are not retained A is back-solved from $5\sigma/S_{\min}$ and is flagged as such in the catalogue, so there the closure is arithmetic and shows it, at 4.0e-07 against 8.7e-05 on the 1641 rows where A comes from the search's own stored geometry. 1641 is therefore the number of rows on which the audit is a test. Nothing else needs correcting, because the detection statistic of Eq. 1, the control ranks and the attribution are all *ratios* formed inside a single window, in which A divides out identically; no crossing, no rank and no disposition depends on it.

Two catalogue columns, `pb_offset_arcsec` and `pb_atten`, carry the offset and the response per window so that a reader can repeat the calculation. They are needed, because the only other geometry column, `theta_pb_arcsec`, is the pipeline's $1.22\lambda/D$ evaluated at $D = 12\text{m}$ for *every* window — including the 1054 rows taken on the $1.7\times$ smaller ACA dishes, where it is correspondingly too narrow — so recomputing A from it alone would give the wrong answer on the majority of the catalogue. The annulus geometry keeps the pipeline's own $\theta_{\text{PB}} = 1.22\lambda/D$, an Airy first-null coefficient; recomputing it on ALMA's FWHM $1.13\lambda/D$ at each window's own dish diameter moves 75 of 1651 windows by more than 1 per cent, extremes -2.6 to $+10.3$ per cent, and no headline threshold moves.

The beam model itself is a bounded systematic, and it is bounded by measurement. `pb_atten` reproduces the Gaussian at $1.22\lambda/D$ that the pipeline applied; CASA's blocked-Airy response is the better physics. Evaluated window by window, the largest ratio between the two over the whole release is 1.0745, and the counts of systems reaching 10^{15}W are identical under either model.

The injection campaign is blind to A by construction, and the resulting bias is exactly zero. The injector deposits an unattenuated amplitude at the star's own direction cosines, which would bias the measured completeness if the ladder were calibrated in true flux — and it is not. On all 40 scored units the ladder is referred to the *apparent* 5σ , so an injected apparent amplitude recovered nine times in ten implies a true flux larger by $1/A$ while P_{trig} is the power of a true $5\sigma/A$, and A cancels identically; 4 of the 40 units have apparent and corrected $S_{\min}$ differing by more than a part in a thousand, so the cancellation is tested and not vacuous. Two limitations remain. Being self-consistent in apparent flux, *the campaign can never falsify the beam model*; and injected configurations span $0.011''$ to $4.98''$ (median $0.033''$, $A \geq 0.9102$) whereas 75 released windows on 5 stars lie further out, reaching $A = 0.552$, so for those the completeness is *interpolated rather than measured*. Their 6 crossings are ratios and unaffected, but a limit quoted for one of those stars rests on an extrapolated curve.

Four rows carry a published noise that no retained product reproduces, and they are declared rather than reconciled: 4 windows of HD 285968 in execution block A002_Xd98580_X3bf6, whose published σ is 0.88–1.32 per cent *higher* than the surviving product's, so the published limits there are weaker and not stronger. With the 10 windows whose products are not retained at all, the declared provenance gap is 14 windows. The two attenuation columns are unaffected: for these rows A comes from a retained product at the same position, where $A \geq 0.999960$, and $S_{\min} = 5\sigma/A$ still closes to $2.6\text{e-}05$.

THE SPATIAL SCREEN: CALIBRATION AND VALIDATION

The spatial screen asks whether an excess sits at the star or anywhere in the field, and it removes most of the unattributed crossings. Everything the dispositions rest on therefore depends on how well that screen is calibrated, which is what this appendix measures. The answer is adverse to the screen and is stated as such: The star-versus-control rank is an empirically useful candidate-prioritisation statistic, but the stellar and control positions are not demonstrably exchangeable under the null. The rank therefore has no calibrated per-window false-alarm probability and is not used as evidence for detection.

The star lies at the phase centre, inside the control annulus, and the asymmetry is real. *It is not a primary-beam effect.* Spectra are extracted from phase-rotated visibilities in apparent flux and are never corrected for the primary-beam response (Appendix A), so the familiar rise of image noise with distance from the pointing centre cannot operate here; and were it operating, it would make the control probes noisier than the star and so place the stellar rank *above* one half, which is the opposite of what is measured. **No mechanism for the displacement is established, and we prefer to say so.** What the control ensemble does show is a weak rise of its own level with radius, and the obvious repair is to remove it: each control probe is standardised within its own radius bin, the probe radii being known from the fixed draw, with no free parameter.

Applied out of sample, the correction does not repair the statistic. It was frozen — probe radii, bin edges and standardisation alike — and applied unmodified to 778 windows that played no part in its construction. The uncorrected stellar rank has median 0.409 against the exchangeable value of 0.5; after correction the median is 0.407, moving further from it, and the result holds separately in both out-of-sample partitions. Rescoring the whole survey with the gradient removed moves the rank-passing list from 12 windows to 11 — losing 2 and gaining 1, all marginal and all dispositioned on other grounds — and leaves the survey median rank slightly further from one half than before.

The interpretation is narrow and, we think, unavoidable. The radial gradient is real, but removing it leaves the median no closer to one half, so the residual is not radial, and neither normalisation yields a calibrated tail probability. **We therefore retain the frozen screen as primary — it has at least the virtue of having been fixed in advance — use it only to decide which windows are examined, and rest every disposition on physical evidence or on recurrence rather than on either rank.** The radial correction, the region-maximum audit and the pseudo-star tests each leave every disposition unchanged.

What the screen costs is not uniform, and almost all of it is paid in a handful of fields. Requiring the star to outrank its ring moves the Class A ninety-per-cent recovery point from $3.06 P_{\text{trig}}$ to $4.42 P_{\text{trig}}$, a factor $\times 1.44$. Where the ring is noise the cost is only $\times 1.15$, screening against a ring peaking at C times the noise being equivalent to raising the trigger to C ; almost the whole remaining effect comes from 4 windows of 402 whose rings carry re-

solved disc emission at $14.1\text{--}18.4\sigma$ (β Pictoris, HD 48370 and HD 139084), where no carrier near the trigger can pass at all.

One mechanism can be excluded by measurement, and it is the one that would matter most. Equation 1 divides every position by a single noise scale, the median absolute deviation across the control probes, shared by the star and all of them. A star bright enough to carry multiplicative calibration residuals of its own would then have a noise that the shared scale does not describe, and its statistic would be inflated with nothing in the search to reveal it. The ratio of the two noises is measurable from the released control vectors alone: because the star and each probe are divided by the same scale and maximised over the same channel $\times$ drift grid, the stellar statistic divided by the median of the 512 control statistics *is* that ratio, and it is one when the two noises agree. Over the windows carrying no threshold crossing its median is ~ 0.44 , with a star-clustered 95 per cent interval of $\sim 0.42\text{--}0.46$. It does not climb with the brightness of the emission in the field, which is what a multiplicative error would scale with: in the fifth of windows with the faintest control rings against in the brightest fifth, and over the windows whose ring exceeds σ and is therefore set by resolved emission rather than by noise. That the ratio can rise at all is shown by the windows that do carry a crossing, where it is ~ 0.44 . Part of the reason it does not rise otherwise is structural: a gain error multiplying a continuum raises every channel of an integration by the same factor, and the per-integration median along frequency removes exactly that. **A stellar-position noise denominator would therefore move no statistic in this survey, and the search has not been re-run with one.**

The reserved hold-out, and what it calibrates

The statistic and the mask were settled on the searched sample and cannot also calibrate themselves, so a hold-out was reserved by a rule fixed and recorded before the archival search finished and before any reserved block was searched: a block is held out if and only if $\text{sha256}(\text{UID})_{[8]} \bmod 5 = 0$, subject to two guards that use identifiers alone — a block on the pre-registered block list deposited with this paper stays in the survey, and no system loses every block — so that the assignment can be recomputed from the UIDs and that list alone. The three choices that can manufacture a candidate were fixed before the rule itself: the statistic of Eq. 1, the promote and retire criteria, and then the reservation, with the first reserved block searched 3 days later. Every diagnostic built afterwards — the radial normalisation, the visibility-domain test, the drift-resolved injection analysis, the P_{90} budget — is incapable of promoting a window.

Of 484 blocks processed, 77 are reserved (15.9 per cent), giving 315 windows toward 36 stars with no star and no system lost from the headline sample. Three quantities are calibrated on those blocks and on nothing else.

First, the rank displacement. Out of sample the stellar rank has median 0.405 where one half is expected, with a block-clustered 95 per cent interval of 0.341–0.458 over 315 windows, and 14 per cent of ranks below 0.1 against 6 per cent above 0.9, so the distribution is displaced rather than growing a tail. The displacement runs the way that

matters: a rank below one half places the star high, and the star in fact outranks all 512 probes in of the survey’s windows against if the positions were interchangeable. The direction therefore inflates star-first outcomes rather than suppressing them, which is why the chance expectations of Appendix D are calibrated on the measured rate. Clustering removes neither the shift nor its size: 99 of 149 blocks have a mean rank below 0.5, and a star-clustered bootstrap places the median at 0.440 (95 per cent 0.42–0.46). The searched sample leans the same way without resolving it (median 0.49, star-clustered 95 per cent 0.47–0.51), and the two distributions are indistinguishable, so the in-sample checks were underpowered rather than contradicted. A displacement that survives a hold-out chosen before the data were seen is a property of the statistic and not of this sample, which is the quantitative reason the rank is a screen and never a significance.

Second, the false-alarm tail factor — the measured rate at which a no-signal position ranks first, divided by the rate exchangeability would give. It is 1.2 on 5 040 trials, with a bootstrap clustered over the 77 reserved blocks giving 0.6–1.9. That factor and its width set the accuracy of every chance expectation in this paper. The reserved blocks contain 0 rank-passing crossings against 0.61 expected.

Third, the radial behaviour of the control ensemble. The standardised control level runs $+0.12, +0.10, +0.12, +0.16$ from the inner to the outer edge of the annulus, a mean per-window slope of $+0.067 \pm 0.033$: the displacement reproduces out of sample, but the steep inner-edge excess seen in sample does not, so the radial term is small, of the wrong sign to account for the displacement, and removing it leaves the displacement behind.

A second out-of-sample set is used here and must not be confused with the hold-out. The *external calibration sample* is 326 execution blocks toward 56 stars, reserved by a rule fixed before the archival search finished; it reaches stars outside the 40 pc work list, so it enters no count of the science sample, and it exists only to calibrate the rank statistic and to supply repeat observations the hold-out does not have. Over its 1322 windows the frozen pipeline returns 7 rank-passing crossings — 4 of them β Pictoris CO, the positive control recovering on data it had never seen — and 3 unattributed against 3.7 expected at the measured tail rate.

One crossing through the whole chain

CP–72 2713 is the clearest worked example of the chain the survey contains, because every stage returns a verdict rather than a shrug. In Band 7 the crossing reaches $T_\star = 5.809$ in channel 35 of its window and outranks every one of the 512 control positions, so it is a candidate; no transition of the frozen mask lies within $\pm 50 \text{ km s}^{-1}$ of it, so it is unattributed; and the drift-following visibility fit puts the emission at the stellar position, $\text{Re}/\sigma = 5.1027$ against a control-position maximum of 1.8222, which by the argument of §5.5.1 is a consistency check and not a promotion.

What disposes of it is the repeat block, and the number quoted for that needs its method attached. The second block matches the first in tuning, channel width, channel count, integration count and on-source time, registers to 0.18 of a channel, and is the deeper of the two, so a

persistent emitter at the first block’s flux would have returned $T_* = 6.21$ where it returns 3.60. Read as two measurements of one constant amplitude — the conservative reading, and the one this paper quotes — the pair is inconsistent at 4.6σ ; taking the discovery flux as exact instead would give 6.8σ . The blocks are 2.05 h apart, so this is a within-night test: it excludes a persistent emitter of that strength on that baseline and nothing about an intermittent one. **The disposition rests on the recurrence test and not on the rank**, which is the chain applied to every crossing this survey cannot attribute.

THE VELOCITY-SPACE LINE MASK, AND WHAT IT COSTS

Attribution decides the candidate list and the mask decides attribution, so the mask’s definition and its price are both given here.

Definition. A one-channel coincidence test is too tight for circumstellar gas: a native channel spans a fraction of a km s^{-1} while real emission is shifted by many channels. The tolerance follows from the kinematic budget of the gas rather than from the instrument — Keplerian motion where these molecules survive ($\lesssim 13 \text{ km s}^{-1}$ beyond 10 au) and linewidths $\lesssim 1 \text{ km s}^{-1}$ — and is $\pm 50 \text{ km s}^{-1}$ around each catalogued transition.

The frame, which is the star’s own. Searched frequencies are topocentric as the correlator delivers them. Before any transition is compared with a crossing, that crossing’s frequency is carried to the barycentre using its own block’s pointing and mid-time and then to the stellar rest frame using the star’s systemic velocity. The barycentric term runs from to km s^{-1} across the survey, so a tolerance applied in the observed frame would spend km s^{-1} of a $\pm 50 \text{ km s}^{-1}$ budget on the Earth’s motion, and the offsets of individual crossings differ between the two frames by up to 55 km s^{-1} . The epochs show the transform working in the only way that can be checked: km s^{-1} apart in the observed frame, they agree to km s^{-1} in the star’s.

The mask holds transitions of species at full JPL/CDMS precision (Pickett et al. 1998; Müller et al. 2005). The species criterion is the one objective rule statable in advance: the transitions ALMA’s band definitions are built around — the CO isotopologues, HCN, HCO^+ , CS, SiO, CN, H_2CO , SO and [C I] — which dominate the emission of disc-selected fields. Sulphur monoxide belongs there for the same reason carbon monosulphide does: of its transitions fall in the intervals searched, it is a routine component of the Band 6 and Band 7 set-ups these fields were observed with, and sulphur bearing molecules are expected wherever the carbon bearing ones already masked are. It is not a formality — the strongest crossing that passes the spatial screen, at $T_* =$, lies km s^{-1} from in its own star’s frame (§5.5).

The half-width is not a tuned parameter, and the counts rather than the argument say so. Table 10 carries the whole disposition chain at half-widths from $\pm$ to $\pm \text{km s}^{-1}$. One count moves, and it is worth naming: narrowing to $\pm \text{km s}^{-1}$ releases from attribution at km s^{-1} from — a crossing that does not outrank its own control ring either way. Everything the result turns on is flat: crossing outranks every control at every half-width in that range, and it is the same crossing each

TABLE 10

THE DISPOSITION CHAIN AGAINST THE MASK HALF-WIDTH, IN THE STELLAR REST FRAME. ALL 56 CROSSINGS ARE CLASSIFIED AT EVERY HALF-WIDTH: THE SECOND COLUMN IS WHAT THE MASK COSTS IN SEARCHED BANDWIDTH AND THE REST IS WHAT IT DECIDES, ENDING IN HOW MANY UNATTRIBUTED CROSSINGS OUTRANK ALL 512 CONTROLS AND HOW MANY OF *those* RECUR AT A LATER EPOCH. ADOPTED HALF-WIDTH IN BOLD.

Half-width (km s^{-1})	band masked (per cent)	line- attributed	unatt- ributed	outrank all controls	recur- ring
± 20	2.1	18	38	1	0
± 25	2.6	19	37	1	0
± 50	4.8	19	37	1	0
± 100	8.7	19	37	1	0

time, while the fraction of searched bandwidth the mask removes rises from to per cent. Widening the tube therefore buys nothing and costs coverage, and narrowing it releases an identified carbon monoxide line of this survey’s own.

The species list is a choice, and widening it is not free. The alternative selection — every Splatalogue transition over the 46 searched islands with $E_u \leq 150 \text{ K}$ and $\log A_{ij} \geq -5$, which is 5,114 transitions of 235 species — masks 75.1 per cent of the union against per cent here, and it does not leave the dispositions alone. of the 56 crossings change side under it, of them a crossing that outranks every control: falls km s^{-1} from a vibrationally excited transition, which nobody would predict toward a quiet main-sequence G star. That is the calibration and not a defect in either list: at these frequencies a uniformly random frequency in one of these windows lies within $\pm 50 \text{ km s}^{-1}$ of something in the wide harvest most of the time. “Coincident with a known molecular line” is a statement about a catalogue rather than a property of a frequency, and a survey quoting one should quote the occupancy of the list it used.

The cost of the mask

Masking removes GHz, per cent of the GHz unique-frequency union, CN, CO and SO dominating. It is not uniform: the loss reaches per cent in Band and stays below 2 per cent in Bands 3–5, and it is tabulated band by band in the data release. The fine-channel experiment pays the larger fractional price, its coverage being concentrated in the line-tuned windows.

What masking buys is classification rather than detection: a carrier at a line frequency is as detectable as anywhere else but cannot be told from the line, because the astrophysical prior at a catalogued transition is overwhelming. The consequences for the limits are in §6.3, and a design that searches every frequency and attaches a *lower* prior to a line-coincident event (Cocconi & Morrison 1959) is in §6.4; the released channels support that second search.

THE FALSE-ALARM EXPECTATION AND THE CONTROL ENSEMBLE

The paper’s central claim is that the number of unattributed crossings is the number chance predicts. This appendix gives the expectation that claim rests on, what it is conditioned on, and what the control ensemble that calibrates it can and cannot support.

What the expectation is conditioned on

The expectation is built from the control statistics themselves, which are the only empirical null the data provide, and it imposes the condition an event must actually satisfy. Under the disposition chain this paper applies, that condition is to reach the trigger *and* outrank every control. Drawing one control index per execution block — so that windows sharing weather, calibration and correlator setup are not treated as independent trials — and resampling blocks, the expectation over the 402 Class A windows is 0.89 unattributed rank-passing crossings against the 2 observed, and the probability of that count or more is $P(\geq 2) = 0.22$. Clustering changes the number little, because the *position* of a block’s control maximum is not shared across its windows.

The conditioning is restrictive, and that is what makes it informative. Exactly one probe per window is the ensemble maximum, so the probability that the stellar position would have been that maximum, treated as exchangeable with the probes, is $1/(512 + 1)$ in windows whose control maximum reaches $T_{\text{trig}} = 5$ and zero in the rest. Only 447 of the 1614 windows that retain a control vector have a control maximum reaching the trigger at all, and 390 of those are Class A. *This is why every rank-passing crossing in the survey is Class A*, and why the choice of reference class largely dissolves once the trigger condition is imposed. An expectation conditioned on the rank alone predicts a different quantity — the number of windows in which a no-signal position would rank first — and is a pipeline diagnostic rather than a statement about candidates.

A second conditioning is reported beside it, on the same control ensemble, and it is the one the body quotes: *without* the rank clause — trigger, localisation and attribution only — 36.5 Class A windows are expected to trigger by chance, the visibility fit localises 0.986 of events at the trigger, and the mask occupies 4 per cent of the searched band, giving *34.5 expected unattributed Class A crossings against 35 observed*. The two are not competing readings of one quantity but the expectations for two different chains, the rank clause being a further one-in-(512 + 1) filter; the observed population is accounted for on either.

A third treatment is quoted for comparison and not as a further expectation: applying the measured non-exchangeability of the stellar position as a factor $\times 1.4$ gives 1.3 expected over the 1614 windows with stored ensembles, 95 per cent interval 0–4. The difference is that factor and nothing else, and it is itself known only to 1.2–1.9; the probability of the observed count or more under it is 0.4 to one significant figure, and no inference is drawn from either. That factor matters because the screen’s rank bias runs against us and not for us: it places the star *high* in its own control distribution (Appendix B), so it inflates star-first outcomes rather than suppressing them. Applying it is the conservative course, and the observed count is still the expected one with it applied.

One trade is not available to us, and it should be stated plainly. An excess of this size needs only a small unmodelled tail in the control ensemble, and Appendix B shows that one exists. But the same noise model sets the injection recovery fractions, so a tail large enough to remove the excess would make the quoted completeness

optimistic by roughly the same factor. The two cannot be traded independently, and neither correction has been applied.

The control ensemble: its calibration, and its resolution

The screen’s geometry is set in §4.1; its calibration and its resolution are the subject here. Under that geometry the stellar ranks are consistent with the uniform distribution exchangeability implies, in both array strata — 597 12m windows and 1048 ACA windows, the latter with median rank 0.478 against the one half expected — over the 1645 of 1651 windows that retain a control vector, the remaining 6 having none. The ACA windows are screened on the geometry measured from each observation’s own antenna table and baselines rather than on a 12m beam, and the array configuration does not enter the sensitivity. Two cautions keep this from being a proof of exchangeability. With samples of this size a shift of less than about 0.042 in the rank distribution could not have been seen at all, so these strata are consistent with uniform rather than demonstrably uniform; and the diagnostic most directly sensitive to contamination of the inner controls, the excess of the innermost control bin over the outer annulus, is +0.038, with an interval of -0.214 to +0.282 that spans zero and therefore neither confirms nor excludes a residual effect.

Neither form is a significance test, and the reason is structural. Ranking the star against 512 control positions drawn from the same field is a powerful way to decide which windows deserve examination, but $1/(512 + 1) = 1.95 \times 10^{-3}$ is the *resolution* of the ensemble — the smallest rank it can express — and not the probability that noise produced the window. Against a survey-wide Bonferroni scale

$$\alpha_{\text{survey}} = \frac{0.05}{1651} = 3.0 \times 10^{-5}, \quad \frac{1.95 \times 10^{-3}}{\alpha_{\text{survey}}} = 64, \quad (\text{D1})$$

that floor is 64 times too coarse, and correlation widens it: neighbouring controls are correlated below the synthesised-beam scale, so a window’s annulus carries only $N_{\text{eff}} \simeq 30\text{--}43$ independent spatial trials and the ensemble resolves no better than 0.023–0.032. Drawing more controls does not rescue it, the primary beam bounding how many independent positions a single pointing contains. **The floor of Eq. D1 is the design’s central limitation: no single window in this survey can be significant on the rank alone, which is why a technosignature claim here requires independent recurrence or independent astrophysical confirmation.** What a visibility-domain statistic buys instead is measurable on this survey’s own data: in the image plane every rank-passing window sits at rank 1 of 512 + 1 with no dynamic range left to separate them, whereas the visibility fit separates the same windows by 5.2σ and its strongest reaches 9.5σ .

The trials budget

A window holds 177–6 869 740 channel $\times$ drift cells (median 468), the survey 1.32×10^8 in total, and they are not independent trials: channel smoothing and a drift grid oversampled two to one together over-count them by 5.3. The control ensembles calibrate that budget without a

Gaussian assumption. The same arithmetic predicts 476 of the 1651 windows carrying a control above T_{trig} and 455 do, agreeing to 4 per cent; corrected for the over-count, the observed control maxima sit +6 and +8 per cent above prediction in the fine and coarse classes independently. Non-Gaussianity of that size is why the operative statement is the empirical rank and not the Gaussian budget.

The residual on-star excess is not zero, and it is reported as an open systematic rather than explained away. Correlation cannot move the expected number above the trigger, only its clumping, so the excess over an iid-Gaussian budget divides in two: the cells within one crossing window, a factor 1.4 that the frequency-axis over-count predicts because one real excursion is counted about three times, and the number of crossing windows, a factor 1.3 with a one-sided Poisson probability of 0.04. The residual runs one way and carries the same sign at the controls. The crossing list is dominated by CO(2→1) at 230.5 GHz toward disc hosts, and a residual excess of chance crossings at the tens-of-per-cent level cannot be excluded.

One stratum departs from uniformity and it is the one that matters. The fine-channel windows' ranks are not uniform, at $p = 0.016$ over $n = 402$ windows, because 12 of them contain an on-star peak above their own control maximum against 0.78 expected; the coarse stratum, three times larger and predicting 2.43 such windows, contains 0. A 15.6–31.25 MHz channel dilutes a narrow circumstellar line, so the class that can resolve one is the class that finds it, and the same reasoning says a genuine unresolved carrier would appear in Class A first. The departure is nonetheless genuine and survives removal of every fine window carrying a stellar crossing, so fine-window rank probabilities are optimistic for the class as a whole and not only where a crossing sits.

What is left is bounded, and the bounds are null. Treating each control as a pseudo-star and ranking it against the other 512-1 gives 220 672 ranks that are uniform to the resolution the ensemble can express, and the real stellar ranks agree with them at $p = 0.38$. Any systematic radial suppression of the controls is below 10 per cent in the mean, and the star's own residual variance equals the controls' to better than a per cent in the rank-passing windows. Scrambling each window so that nothing coherent survives along a drift track carries 4 windows past the rank floor and changes no disposition; it fails outright only on HD 48370 Band 6, whose extended emission and imaging residuals are coherent across integrations and are destroyed by the scramble itself, so no scrambled-noise probability is quoted there. The remaining stratifications — by per-control noise and by gain — cannot be made at all, because neither is kept in the retained products.

RADIO-FREQUENCY INTERFERENCE SCREENING

Interference is the first alternative explanation for an unattributed crossing, so the screening this survey has, and the coverage it does not have, are both set out here.

Millimetre interferometry is an unusually poor environment for terrestrial interference, though not an empty one: ITU Radio Regulations Article 5 carries active allocations throughout 90–275 GHz, and WRC-19 identified

parts of 275–450 GHz for land-mobile and fixed use. The specific hazard is the 94.0–94.1 GHz Earth exploration-satellite allocation used by spaceborne cloud radar — CloudSat and, since 2024, EarthCARE — a documented contaminant of ALMA Band 3; 1 of the 1651 searched windows overlaps it, carrying 0 threshold crossings and 0 rank-passing events, so no result here rests on it. What favours this band is that terrestrial transmitters are sparse compared with centimetre wavelengths, that the array sits at 5000 m on a site chosen partly for radio quietness, and, decisively, that an interferometer is not a total-power detector: a signal not in the far field at the phase centre does not fringe-track but acquires a rapidly varying geometric phase across the array, and is suppressed in the correlation rather than imaged.

What screening exists, stated with its coverage, because the coverage is uneven. The allocation test above and the sky- and intermediate-frequency tests below are computed for every searched window from the released catalogue and are reproducible from it, as is the 512-position rank screen, which asks whether an excess is at the star or everywhere in the field and so serves as an interference test as well as a prioritisation. Recurrence covers only the crossings that have a covering repeat block, and the visibility fit — which asks whether an excess has the phase signature of a source at the phase centre, as a near-field emitter would not — is a consistency check rather than a filter. The correlator flagging is the observatory's own, replayed verbatim from each block's delivered calibration script, but it is not recorded per window in the release. *No per-channel interference flag of our own exists, and none can be recovered.*

Four measurements support the reasoning, and each is a test the population could have failed.

First, fixed sky frequency. Interference is a property of the observatory and not of the star, so a terrestrial carrier should appear at the same sky frequency in every window that covers it. Of the 2 unattributed rank-passing events, 2 have a released crossing frequency; 32 windows toward 8–24 other stars in 74 other blocks cover those frequencies, and 0 of them carries a crossing in the same channel.

Second, clustering in sky frequency. If interference dominated the crossing population, the crossings would pile up at particular sky frequencies. They do cluster — but on the survey's own molecular lines. Of the 51 crossings with a released frequency, 11 fall inside the mask's own $\pm 50 \text{ km s}^{-1}$ velocity window against 4.7 expected if each window's peak fell anywhere on that window's own channel grid, $p = 0.0007$. The densest band, 230.423–230.723 GHz, holds 12 crossings, of which 9 are within that velocity window of CO(2→1); and once the line-attributed crossings are set aside, the 40 that remain show no clustering at all ($p = 0.111$). *The concentration is the attribution, not a residue left over after it.*

Third, fixed intermediate frequency. An artefact of the signal chain would sit at a fixed position within the spectral window rather than at a fixed sky frequency. This test is run over every window for which the catalogue releases a peak frequency — of them, of which only carry a threshold crossing — because such an artefact would occupy the same intermediate frequency in all of them and not only where the trigger fires. It does not: the peaks' fractional positions within their own windows

have median 0.50 and are uniform.

Fourth, a cross-target occupancy screen over every searched window and not only the crossings. A terrestrial carrier should recur at the same sky frequency toward *unrelated* stars. Over all 3027 searched windows toward 133 stars, with each window’s peak randomised on its own channel grid, the number of cross-star coincidences within 5 MHz is 569 against 554 ± 13 expected ($p = 0.15$); among the 109 windows whose peak exceeds the trigger it is 4 against 5.2 ± 1.9 , i.e. *below* the null, and with the molecular-line mask applied 1 against 2.1 ± 1.3 . There is no excess to explain. The screen is keyed on the star rather than on the product directory, which matters: a star’s own repeat observations are not different targets, and every persistent feature — a real line most of all — would otherwise enter as a cross-target coincidence. 35 of those 109 windows sit on a masked transition (CO(2–1) $\times$ 17, CO(1–0) $\times$ 10, CO(3–2) $\times$ 4, ^{13}CO (2–1) $\times$ 3, SiO(5–4) $\times$ 1), which is the point: a CO line falls at almost the same sky frequency toward every nearby star, so cross-star frequency coincidence is exactly what real line emission produces, and the screen cannot distinguish it from interference unless the mask is applied first.

One execution block fails these tests, and it fails them as a block rather than as a carrier. Its crossings toward BD+05 1668 are the highest statistics in the survey outside β Pictoris and HD 48370, and the reason is in the field and not at the star: all 512 control positions exceed the trigger, the ring’s own median running – against a survey median of , so the window contains no null against which anything could be localised. The visibility fit returns a large imaginary part, placing the emission away

from the phase centre; of the sit at the identical channel of two different spectral windows, which is the signature of a fixed intermediate frequency; and the block’s 24 sibling executions of the same field, with the same array and the same spectral setup, produce no crossing at all. It is *not* a terrestrial carrier at a fixed sky frequency: each of the four frequencies is covered by – further windows toward – other stars, and of them carries a crossing in the same channel. **The evidence supports a data-quality failure of one execution block, and that is how these four crossings are dispositioned (§5.3.2):** under a rule fixed before it was applied — a window is flagged when the median of its own control ensemble exceeds the survey median of that quantity — of the 1651 windows are flagged, toward stars (), the remaining one being a resolved disc rather than a defect. The flag is a statement about the window, not an attribution of the feature.

None of this is a substitute for the ON/OFF cadence a dedicated search uses, and we do not present it as one. Accordingly we claim only that *no identified interference mechanism accounts for the unattributed events*: allocation tables cannot exclude harmonics, intermodulation products, unconventional emitters or near-field sources, and an archival dataset without a simultaneous off-source strategy cannot do better. What the four tests establish is that the crossing population as a whole carries none of the signatures interference leaves — fixed sky frequency, fixed intermediate frequency, or presence in unrelated targets — which is a statement about the population and not a clearance of any individual crossing.

REFERENCES

- Astudillo-Defru N., et al., 2017, *A&A*, 602, A88
 Astropy Collab., 2022, *ApJ*, 935, 167
 Backus P. R., et al., 2002, *BAAS*, 34, 1173
 Blomme R., et al., 2023, *A&A*, 674, A7
 CASA Team, 2022, *PASP*, 134, 114501
 Burton K., MacGregor M. A., Osten R. A., 2022, *ApJL*, 939, L6
 Burton K., 2024, PhD thesis, Univ. Colorado Boulder, ch. 4
 Cataldi G., et al., 2023, *ApJ*, 951, 111
 Choza C., et al., 2024, *AJ*, 168, 254
 Cocconi G., Morrison P., 1959, *Nature*, 184, 844
 Czech D. J., et al., 2021, *PASP*, 133, 064502
 Dent W. R. F., et al., 2014, *Science*, 343, 1490
 Drake F. D., 1961, *Physics Today*, 14(4), 40
 Drake F. D., 1974, in Sagan C., ed., *CETI*. MIT Press
 Enriquez J. E., et al., 2017, *ApJ*, 849, 104
 Gaia Collab., Vallenari A., et al., 2023, *A&A*, 674, A1
 Garrett M. A., et al., 2026, *MNRAS*, in press (arXiv:2605.10212)
 Gentile Fusillo N. P., et al., 2021, *MNRAS*, 508, 3877
 Ginsburg A., et al., 2019, *AJ*, 157, 98
 Harris C. R., et al., 2020, *Nature*, 585, 357
 Huang B.-L., et al., 2026, *ApJ*, 1006, 9
 Jacobson-Bell K., et al., 2025, *AJ*, 169, 233
 Johnson O. A., et al., 2023, *AJ*, 166, 193
 Li J.-K., et al., 2026, arXiv:2603.19023 (preprint)
 Ma P. X., et al., 2023, *NatAs*, 7, 492
 MacGregor M. A., et al., 2020, *ApJ*, 891, 80
 Manunza L., et al., 2025, *AcAau*, 233, 155
 Margot J.-L., et al., 2023, *AJ*, 166, 206
 Mason L. A., Garrett M. A., Wandia K., Siemion A. P. V., 2025, *MNRAS*, 536, 2127 (arXiv:2411.19827)
 Matrà L., et al., 2017, *MNRAS*, 464, 1415
 Morrison I. S., et al., 2023, *PASA*, 40, e038
 Müller H. S. P., et al., 2005, *JMoSt*, 742, 215
 Ortiz Ceballos K., et al., 2026, arXiv:2609.16720 (preprint)
 Pickett H. M., et al., 1998, *JQSRT*, 60, 883
 Poznanski D., et al., 2025, *AJ*, in press (arXiv:2505.03927)
 Price D. C., et al., 2020, *AJ*, 159, 86
 Sheikh S. Z., et al., 2025, *AJ*, 169, 118
 Sheikh S. Z., 2020, *IJAsB*, 19, 237
 Sheikh S. Z., et al., 2021, *Nature Astronomy*, 5, 1153
 SpaceX, 2024, FCC IBFS SAT-MOD-20230207-00021
 Tarter J., 2001, *ARA&A*, 39, 511
 Tremblay C. D., et al., 2024, *PASA*, 41, e004
 Vidal C., et al., 2026, arXiv:2605.21093 (preprint)
 Virtanen P., et al., 2020, *NatMe*, 17, 261
 White G. J., et al., 2026, *MNRAS*, submitted
 Wright J. T., et al., 2018, *AJ*, 156, 260
 Zhang Z.-S., et al., 2020, *ApJ*, 891, 174